WHITEPAPER

the agentic HR operating model

what changes when intelligence acts inside the process rather than reporting on it, with a maturity model

7 chapters · 35 sources · entomo research

entomo

2026

CONTENTS

what is in this whitepaper

EXECUTIVE SUMMARY

the argument in short

Most HR technology of the last fifteen years reported on a process after it ran. Agentic systems change the state of a process while it is running. A dashboard flagging elevated attrition risk has changed nothing. A system that opens a retention review, assigns a named manager, drafts the conversation guide and sets a follow-up date has moved the process. That is a difference of authority, not sophistication, and it rewrites who is accountable and what evidence survives.

The shift is being attempted widely and landing rarely. SHRM's 2026 study of more than 1,900 HR professionals found 39% of organisations have implemented AI anywhere in HR, concentrated in recruiting at 27%, HR technology at 21% and learning at 17%. A Gartner survey of 114 HR leaders found 88% saying their organisations have not realised significant business value from AI tools. MIT's Project NANDA reported roughly 95% of generative AI pilots producing no measurable profit and loss impact. The common thread is not model quality. The tools sat beside the workflow instead of inside it, and nobody defined the decision they were meant to change.

This paper offers two named instruments you can run without buying anything. The Agentic HR Maturity Ladder places your function on one of five stages by what the system may do and what evidence you can produce, not by which vendors you have signed. The Reversibility and Reach Test scores any HR task in two minutes and assigns it to one of three authority rungs: propose, decide, or execute. Three supporting instruments follow: the Meaningful Review Test for a checkpoint that is not decorative, the Eight-Field Decision Record for the log that has to survive discovery, and the 30/30/30 Pilot with four numeric gates.

The regulatory clock is less forgiving than the delayed headlines suggest. The EU pushed standalone high-risk obligations under Annex III of the AI Act from August 2026 to 2 December 2027 through Regulation (EU) 2026/1744, published on 24 July 2026. That delay does not travel. California's automated decision system rules took effect on 1 October 2025 and require four years of retention. Illinois amended its Human Rights Act from 1 January 2026. New York City has required annual independent bias audits since 2023. India notified its Digital Personal Data Protection Rules on 13 November 2025. Singapore's PDPC issued advisory guidelines on personal data in AI decision systems in March 2024. Saudi Arabia's PDPL has been fully enforceable since 14 September 2024. All of them converge on three things: a named human, a stated basis, a durable record.

This paper is published by entomo. Every number is attributed to a named public source with a working link, so you can check the reasoning rather than take it. Where we estimate, we say so and show the arithmetic. If you read one section, read chapter five. The checkpoint is where most agentic HR programmes will fail, quietly, and pass their own audit while doing it.

CHAPTER 01

what changes when intelligence acts inside the process

Reporting systems describe a process after it runs. Agentic systems change its state while it runs. Everything else in this paper follows from that.

An analytics layer answers questions. An agentic layer changes records. That is the whole distinction, and it is not a technical one. When a dashboard reports elevated attrition risk, nothing has happened. When a system opens a retention review, assigns a named manager and sets a follow-up date, the organisation holds a new state, a new owner and a new obligation. Nobody has to read a report for the consequence to arrive.

the three things that change

1. **Latency.** The gap between a signal appearing and someone acting on it falls from days to seconds. That is the benefit and the risk, because a wrong action now propagates before anyone opens a report.
2. **Authority.** Someone has decided a system may change a record without asking. That decision is usually made inside a configuration screen, by a person who does not think of themselves as delegating authority, and it is rarely written where a general counsel would find it.
3. **Evidence.** A report can be reconstructed from source data at any time. An action cannot. If a system recommended, a human accepted in four seconds, and the model has since been retrained, the only surviving evidence is the log you chose to write.

HR bought analytics for a decade and a half and judged it on insight quality. That habit does not transfer. An agentic layer has to be judged on decision quality and record quality, because those are the two things examined when a decision is challenged.

39% organisations that have implemented AI anywhere in the HR function

SHRM, 2026

The placement is more revealing than the number. Recruiting sits at 27%, HR technology at 21%, learning at 17%, while inclusion, ESG and compliance sit at 2% or below. HR has adopted AI where the output is text and the stakes are diffuse, and avoided it where the output is a decision about a named person. Agentic systems remove that hiding place, because the output is the decision.

88% HR leaders saying their organisation has not realised significant business value from AI tools

Gartner, 2025

The reason for that gap is documented and it is not intelligence. MIT's Project NANDA analysed 52 executive interviews, 153 leader surveys and 300 public deployments, and found roughly 95% of generative AI pilots produced no measurable profit and loss impact. Deloitte's 2026 Global Human Capital Trends, run with Oxford Economics across more than 9,000 leaders in 89 countries, found organisations taking a technology-first approach were 1.6 times more likely to miss their return targets. Both point at the same failure: the tool never entered the workflow it was bought to change.

> **THE WORKFLOW TEST**
>
> Before funding anything, ask one question. Which record does this change, and who is accountable when it changes wrongly? If nobody can name the record, you are buying a report. If nobody can name the person, you are buying a liability.

where latency already costs money

Payroll is the clearest case, because the errors are countable. An EY survey of 508 respondents, mostly HR professionals at United States companies with 250 to 10,000 employees, found an average of 15 corrections per pay period, 80% payroll accuracy and an average cost of 291 dollars per error. One in six companies faced litigation, averaging 32 complaints and 29 hours of internal time each.

The arithmetic that follows is ours, not EY's. On a monthly cycle, 15 corrections across 12 periods is 180 a year. At 291 dollars each that is 52,380 dollars, before manager hours or goodwill. A system that catches the error before the run is attacking a real number. A system that silently corrects net pay without notifying the employee is creating a different one.

FOUR HR PROCESSES WHERE LATENCY HAS A PUBLISHED PRICE, AND WHAT CHANGES WHEN A SYSTEM ACTS INSIDE THEM

process	published benchmark	source	what acting inside changes	what it puts at risk
payroll correction	15 corrections per pay period; 80% accuracy; 291 dollars per error	EY, 2022	errors surface before the run, not after	a silent correction that alters net pay with no notice
requisition load	median 39 days to fill non-executive roles, down from 44; 67% rise in requisitions per recruiter at extra-large employers	SHRM, 2026	screening and follow-up run continuously	a screening rule that thins a protected group before any human sees the pool
skills currency	39% of skill sets expected to be transformed or outdated between 2025 and 2030	World Economic Forum, 2025	profiles update from work evidence, not annual surveys	an inferred profile that quietly gates internal opportunity
HR capacity	median 1.98 HR staff per 100 employees, against 1.58 in 2017	SHRM, 2025	routine case volume absorbed without headcount	reviewer capacity nobody budgeted, because checkpoints were never costed

Notice the shape of the risk column. In every case the failure is not that the system is stupid. It is that the system is fast, plausible and unobserved. Speed without a checkpoint turns a small error rate into a large error count, and the count is what appears in a complaint.

> **The question is no longer whether the model is good. It is whether anyone can say who decided, on what evidence, and how it could have been stopped.**

One thing does not change. Employment law does not distinguish between a decision made by a person, a policy, a spreadsheet rule or a model. Automating the mechanism does not move the duty, and in several jurisdictions it adds one.

CHAPTER 02

the five stages of agentic HR maturity

Most maturity models grade the technology. This one grades authority and evidence, because those are what a regulator or a claimant will test.

Maturity models fail HR for a consistent reason. They measure what has been bought, so a function with a large licence estate scores well while a function with three tightly governed automations scores badly. The useful question is not how much intelligence you have deployed. It is how much authority you have delegated, and whether you can prove what happened when you did.

You are at a stage below only if you can produce the evidence named at that stage, on demand, for a decision made ninety days ago. That test is what makes the ladder usable. Without it, every function self-reports at stage four.

the Agentic HR Maturity Ladder

Five stages defined by what the system may do and what record survives, each with an evidence test.

01 **stage 1: recorded**

The system stores what happened; reporting is retrospective. Humans do all the sensing, deciding and doing. Evidence test: you can reproduce last quarter's headcount report from source data. Most enterprises cleared this a decade ago and mistake it for analytics maturity.

02 **stage 2: surfaced**

The system ranks, scores, flags and alerts. It has an opinion but no hands. Evidence test: for any score shown to a manager, you can state which input fields produced it and when they were last refreshed. Most functions describing themselves as advanced are here.

03 **stage 3: drafted**

The system proposes a specific action with a rationale, prepared for a human to accept, edit or reject. Nothing moves until a person acts. Evidence test: every proposal carries an identified decider, a timestamp and a reason code, and you can report override rate by process and by reviewer.

04 **stage 4: bounded**

The system executes inside written limits and logs what it did: rescheduling interviews, closing duplicate cases, correcting a time entry below a threshold. Evidence test: the limits exist as a signed document, breaches alert a named owner, and a quarterly sample of executed actions has been reviewed with findings recorded.

05 **stage 5: governed**

Agentic behaviour spans linked processes and the control environment is standing rather than project-based: an agent inventory, an owner each, a tested stop, versioned models, scheduled bias testing. Evidence test is adversarial: outside counsel can reconstruct a contested decision from your records alone, without interviewing anyone who was there.

the thresholds between stages

WHAT HAS TO BE TRUE TO CLAIM EACH STAGE, AND WHAT STOPS YOU MOVING UP

stage	the human's job	evidence you must produce	the exit criterion
1 recorded	interpret periodic reports	reproducible reports from source data	a defined signal someone is accountable for watching
2 surfaced	judge the flag, decide the action	input lineage and refresh date for any score shown	a written statement of which actions the system may draft
3 drafted	accept, edit or reject each proposal	named decider, timestamp, reason code, override rate	override rate measured for two consecutive quarters, inside a stated band
4 bounded	set limits, review exceptions, sample execution	signed limit document, breach alerts, quarterly sample review	an agent inventory with a named owner and tested stop for every entry
5 governed	own the control environment and its exceptions	reconstruction of a contested decision from records alone	continuous, since stage five is maintained rather than reached

Two features matter more than the labels. The work at each stage differs in kind, not degree: a stage three reviewer judges individual cases, a stage five owner judges a control system. Putting the first person in the second seat is a common and expensive substitution. And stages cannot be skipped safely, because each produces the evidence the next depends on. Bounded execution without a measured override rate means you set your limits by guesswork, and you will learn the right threshold from an incident.

15% IT application leaders considering, piloting or deploying fully autonomous AI agents, against 75% doing so for supervised agents

Gartner, 2025

That survey covered 360 IT application leaders at organisations of at least 250 employees across North America, Europe and Asia Pacific, fielded in May and June 2025. Only 19% held high or complete trust in their vendor's protection against hallucination. The market is mostly at stages two and three and correctly reluctant to move further. The gap between 75% and 15% is not timidity. It is the missing control environment that stage four requires.

THE HONEST SELF-ASSESSMENT

Ask your HR technology lead which stage you are at. Then ask a member of your employee relations team to produce the full record of one contested decision from last quarter, unaided, in one working day. The second answer is your real stage.

Regional context changes the pace, not the sequence. PwC's Middle East Workforce Hopes and Fears Survey 2025, covering 1,286 employees, found 75% had used AI tools at work in the previous year against 69% globally. Deloitte's 2026 State of AI in the Enterprise reporting for India found 40% of Indian respondents at significant or full AI usage against roughly 28% globally, with 39% naming regulatory compliance as their top integration challenge. Adoption in these markets runs ahead of the global average while the governance layer is still being built.

place your function on the ladder in one working session

- [ ] List every HR process where software writes to a record without a person pressing something. If the count is zero, you are at stage two or below regardless of what you have bought.
- [ ] Name the individual accountable for each one. Not the team, not the vendor, not the steering committee. If you cannot name a person in thirty seconds, that process is ungoverned.
- [ ] Pick one decision from the last ninety days an employee could plausibly contest, and try to assemble the full record in one day from stored artefacts only.
- [ ] Measure override rate for any process at stage three. If it has never been measured, you are at stage two with a stage three interface.
- [ ] Ask whether a stop exists for each automated process and whether it was tested in the last six months. An untested stop is an intention, not a control.
- [ ] Write your stage on one line, with the single missing piece of evidence that keeps you from the next one. That line is your next two quarters.

Most large enterprises in Southeast Asia, the Middle East and India will land at stage two with pockets of stage three in recruiting. That is an accurate result, and it makes the next question answerable: for which specific tasks should the system move up a rung?

CHAPTER 03

the authority ladder: propose, decide, execute

Three rungs, one scoring test, and one rule that caps the whole thing. This is the shortest chapter and the one to circulate first.

Delegating authority to software is an act of governance, not configuration. It deserves the discipline you apply to signing limits in finance. Nobody would let a system move money without a written limit and a named approver. HR routinely lets systems move people without either.

the three rungs

- **Propose.** The system assembles evidence and sets out options without recommending one. A human constructs the decision. Used where the decision is contested, irreversible or legally named.
- **Decide.** The system recommends one action with a rationale and a stated confidence. A named human accepts, modifies or rejects, and the choice is logged with a reason code. Nothing executes without that act.
- **Execute.** The system acts inside written limits and logs what it did. Humans set the limits, receive breach alerts and sample the output. Used only where the action is cheap to reverse and narrow in reach.

These are not a maturity sequence. A mature function runs all three at once, on different tasks, deliberately. The failure mode is not being on the wrong rung. It is not knowing which rung a given task is on, which is the normal state of most HR technology estates.

the Reversibility and Reach Test

A two-minute test that assigns any HR task to a rung, using two scores from 1 to 5 that add to a placement between 2 and 10.

01 **score reversibility, 1 to 5**

How hard is this to undo? Score 1 if it reverses within one business day at no cost or visibility to the employee. Score 3 if reversal takes a week, involves a second party, or the employee will notice. Score 5 if it cannot be undone: a rejection already received, a message that has left the building, a payment made, a record disclosed.

02 **score reach, 1 to 5**

How many people does one action touch, and how personal is it? Score 1 for one person on an administrative matter. Score 3 for one person on a matter affecting their work, or many people on an administrative matter. Score 5 for a decision about a named person's employment, pay, progression or standing, or any action applying one rule to a whole population at once.

03 **add and place**

A total of 2 to 4 places the task on execute. A total of 5 to 7 places it on decide, with a mandatory checkpoint. A total of 8 to 10 places it on propose, with no default acceptance path in the interface. These bands are our design heuristic, chosen so any score of 5 on either dimension forces at least the decide rung.

04 **apply the cap**

Regardless of score, no task touching hiring, promotion, pay, discipline or termination may sit above propose. These are the categories named law already covers, and the cost of being wrong is not measured in cycle time.

05 **write it down and date it**

Record both scores, the placement, the person who set it and the review date. A placement without an owner and a date is an opinion. Re-score any task whose model, data source or population changes, and re-score everything annually.

THE REVERSIBILITY AND REACH TEST APPLIED TO TWELVE COMMON HR TASKS

task	reversibility	reach	total	rung
reschedule an interview after a declined slot	1	1	2	execute
close a duplicate HR service case	1	1	2	execute
send a benefits enrolment reminder	1	2	3	execute
correct a missing time punch below a stated value	2	1	3	execute
draft a job description from a role profile	1	3	4	execute
flag a payroll anomaly to a named analyst	2	3	5	decide
propose a shortlist order for a hiring manager	3	3	6	decide, capped at propose
recommend a learning path to an employee	2	4	6	decide
assign an employee to a redeployment pool	4	4	8	propose
adjust a performance rating distribution	4	5	9	propose
reject a candidate without human review	5	5	10	propose
initiate a disciplinary process	5	5	10	propose

Two rows deserve attention. Drafting a job description scores 4 and lands on execute, which surprises people, because it feels consequential. It is cheap to reverse and touches no individual record, so the rung is right. Shortlist ordering scores 6 and would land on decide, but the cap pulls it to propose, because ordering a shortlist is selection by another name and selection is what the law names.

WHY THE CAP EXISTS

Annex III of the EU AI Act classes employment and worker management systems as high risk. GDPR Article 22 gives individuals the right not to be subject to solely automated decisions with legal or similarly significant effects. New York City has required annual independent bias audits since 2023. Illinois added notice duties from 1 January 2026. California's rules took effect on 1 October 2025. The cap is not caution. It is the shape of the law.

The most common way the ladder is defeated is the default accept: a recommendation with the accept button pre-focused, a reject path requiring free text, and a queue counter in the corner. That design produces a decide rung on paper and an execute rung in practice. Chapter five deals with it at length, because it is the largest single source of governance theatre in agentic HR.

writing the authority statement

1. **Scope.** The exact task, the population it applies to, and the systems of record it may write to.
2. **Rung and score.** Both scores, the total, the placement, and any cap applied.
3. **Limits.** For execute tasks, the numeric boundaries: value thresholds, volume per day, populations excluded, hours of operation.
4. **People.** The accountable owner, the reviewers if the rung is decide, and the escalation path with a name at the end of it.
5. **Stop and review.** How the process is halted, who may halt it without approval, when the stop was last tested, and the next review date.

8% HR leaders who believe their managers have the skills to use AI effectively today

Gartner, 2025

That figure comes from a Gartner survey of 114 HR leaders in July 2025. A companion survey in August 2025 found only 14% of organisations supporting managers on integrating generative AI into daily work. Both matter here, because managers are the reviewers on most decide-rung tasks. Placing a task on decide assumes a competent decider. If 8% is your honest read, some decide tasks belong on propose until training closes the gap.

> **A recommendation with a pre-selected accept button is not a decision point. It is an execution with extra steps and a signature you can point at.**

Run the test on your twenty highest-volume HR tasks this month. It takes about two hours with three people in a room. The output is a one-page register telling your auditors, your works council and your general counsel exactly what your systems may do. Very few enterprises can produce that page today.

CHAPTER 04

a role-by-role impact map for the HR function

Agentic systems do not remove HR roles evenly. They move work between rungs, and each move creates a specific new failure mode.

The usual framing asks which HR roles are at risk. That produces anxiety and no decisions. A better question is which part of each role moves to the system, what stays irreducibly human, what new failure that creates, and which single metric catches it early. The map below answers those four questions for nine roles.

WHAT MOVES, WHAT STAYS, AND THE METRIC THAT CATCHES THE NEW FAILURE

role	what moves to the system	what stays human	the new failure mode	the metric that catches it
CHRO	standing reporting, scenario modelling, first-pass workforce plans	setting the authority ladder, owning the control environment, defending decisions externally	governing by dashboard, where the summary is trusted because it is fast	share of agent-executed processes with a named owner and a tested stop
HR business partner	meeting preparation, case summarisation, drafting manager guidance	reading a room, judging a manager's motive, deciding what not to escalate	becoming a routing layer for system-generated actions rather than an adviser	share of HRBP time in scheduled advisory work versus queue clearance
talent acquisition lead	sourcing, scheduling, candidate messaging, interview scaffolding, requisition triage	final selection, offer judgement, hiring manager calibration	a screening rule that thins a protected group before any human sees the pool	pass-through rate by demographic group at each funnel stage, monthly
people analytics lead	report production, data preparation, anomaly detection, routine queries	defining what counts as evidence, challenging inference, setting standards	the function becomes a model steward and loses independence to challenge	agent recommendations formally contested by analytics per quarter
HR operations and shared services	tier zero and tier one resolution, case routing, document generation	exception handling, distress detection, judgement on ambiguous eligibility	silent auto-resolution of cases the employee did not consider resolved	reopen rate within fourteen days, split by auto-resolved and human-resolved
total rewards	market data assembly, range modelling, letter generation, anomaly flagging	pay decisions, equity judgement, defending outcomes to works councils	a recommendation engine encoding last year's inequity as this year's benchmark	pay gap movement where recommendations were accepted unmodified

role	what moves to the system	what stays human	the new failure mode	the metric that catches it
learning and development	content assembly, skill inference, path recommendation, nudging	deciding what capability the organisation needs and will not build	inferred skill profiles that quietly gate internal opportunity	share of employees who have viewed and corrected their inferred profile
employee relations	case intake, precedent retrieval, documentation and timeline assembly	investigation, credibility assessment, sanction, every contested conversation	precedent retrieval that narrows the frame before the investigator forms a view	share of cases where the investigator recorded a line of inquiry the system missed
HRIS and HR technology	integration monitoring, configuration testing, access review, log collection	agent inventory ownership, version control, stop design and testing	an estate of agents nobody has enumerated since the last renewal	production agents in the inventory versus agents found in an audit

the three roles that change most

Talent acquisition changes first, because volume makes automation attractive and the law makes it dangerous. SHRM's 2026 benchmarking, from more than 4,600 organisations, put the median time to fill a non-executive role at 39 calendar days, down from 44, while extra-large employers reported a 67% rise in median requisitions per recruiter. Take a recruiter carrying 15 open requisitions. Our arithmetic on that published change: 15 multiplied by 1.67 gives 25.05, so roughly 25.

39 days median time to fill a non-executive role, down from 44 days the previous year

SHRM, 2026

The recruiter carrying 25 requisitions instead of 15 has a real problem, and an agent solves part of it. The ladder forces the question of which part. Scheduling and follow-up score low on both dimensions and belong on execute. Shortlist ordering is capped at propose. A function that automates both without distinguishing them has not added capacity. It has moved a selection decision into a queue nobody has time to read.

People analytics changes in a way that is easy to miss. When the function producing reports becomes the function operating the models, its independence to challenge them erodes. The team that built the attrition score is now the team asked whether the attrition score is fair. Keep the challenge function separate from the build function, even if it is one person with a different reporting line for two days a month.

Employee relations changes least, which is correct. Investigation, credibility assessment and sanction sit at the top of the test and stay on propose permanently. What changes is preparation, where precedent retrieval and timeline construction are genuine savings. The risk there is subtler than bias: a system that surfaces three similar past cases frames the investigator's thinking before they have formed a view.

> **THE MANAGER PROBLEM SITS UNDERNEATH ALL OF THIS**
>
> Nearly every decide-rung task in this map is reviewed by a line manager, not by HR. Gartner found only 8% of HR leaders believed their managers could use AI effectively, and only 14% of organisations supporting managers on integrating it into daily work. You cannot design a checkpoint for a reviewer you have not equipped. Budget the enablement before the deployment, or move the task down a rung.

the two roles this creates

- **Agent owner.** A named individual accountable for one agent's behaviour, limits and outcomes. Not the vendor, not the project. They approve the authority statement, receive breach alerts, hold the stop and sign the quarterly sample review. Our working threshold: no more than five production agents per owner, because ownership past that becomes nominal.
- **Checkpoint owner.** Accountable for whether human review on decide-rung processes is real. They measure override rates, watch reviewer capacity, test interfaces for default-accept patterns, and may move a task down a rung without asking permission. This role should not report to the person whose cycle time the agent is improving.

Both roles are part-time for most enterprises and both are usually left unassigned. That omission is the most reliable predictor of a governance failure, because every control in the next two chapters depends on somebody being accountable for running it.

On headcount, be honest in both directions. SHRM's 2025 CHRO benchmarking put the median HR-to-employee ratio at 1.98 per 100 employees against 1.58 in 2017, with HR expense at 2.4% of operating expense and median budgets rising 9.1% year on year. HR grew through the analytics era. Agentic systems will reduce transactional load and add review, monitoring and record-keeping load that did not exist before. Model both, or the business case will be reopened within a year.

> **Automation moves work. It rarely removes it. The functions that are surprised are the ones that budgeted for the saving and not for the supervision.**

Mercer's Global Talent Trends 2026, from nearly 12,000 respondents surveyed in September and October 2025, found 65% of executives expecting between 11% and 30% of their workforce to be redeployed or reskilled because of AI, and 98% planning organisational design changes within two years. Meanwhile 51% of C-suite respondents felt prepared for the human and machine era, down from 65% in 2024. Confidence is falling as intent rises, which is a reasonable response to a real difficulty.

CHAPTER 05

designing the human checkpoint

Human review is the control everyone claims and few can evidence. The research says adding a human can raise adoption while lowering accuracy.

The phrase human in the loop has become a compliance answer rather than a design specification. It appears in vendor documentation, board papers and regulatory filings, almost never with a definition of what the human must be able to do. The experimental evidence is uncomfortable and worth knowing before you design anything.

66% of the time participants chose to delegate a prediction to an algorithm over an equally accurate human

Sele and Chugunova, PLOS ONE, 2024

That study ran an online experiment with 292 participants on a prediction task. Preference for the algorithm rose a further 7 percentage points when participants could monitor and adjust its output. Crucially, participants were *less* likely to intervene on the least accurate recommendations. The authors' conclusion is the sentence to carry into your next design review: the human-in-the-loop design increased uptake and decreased accuracy.

The clinical literature is older and larger. Goddard, Roudsari and Wyatt screened 13,821 papers and included 74 in a 2012 systematic review of automation bias in the Journal of the American Medical Informatics Association. Related work by Goddard on prescribing decision support measured the rate directly: 5.2% of cases involved a clinician switching from a correct pre-advice decision to an incorrect post-advice one, at 70% system accuracy, across 26 general practitioners and 520 cases. The same system lifted net accuracy from 50.4% to 58.3%. Both facts are true at once, and that is the point.

65.8% increase in commission errors when decision support gave incorrect advice on low-complexity tasks

Lyell and colleagues, BMC Medical Informatics and Decision Making, 2017

That 2017 study of 120 final-year medical students found correct decision support reduced omission errors by 38.3% to 46.6% depending on conditions, while incorrect support increased omission errors by 24.5% to 33.3% and commission errors by 51.7% to 65.8%. Read the asymmetry carefully. Good advice helps by a lot. Bad advice hurts by more. A checkpoint that does not change behaviour when the recommendation is wrong is not a control, whatever the process document says.

The EU AI Act reaches the same conclusion in legal language. Article 14(4)(b) requires that people assigned to oversight remain aware of the tendency to rely automatically or over-rely on a high-risk system's output, naming automation bias explicitly. Article 14(4) also requires that they can understand the system's capacities and limits, correctly interpret its output, disregard or reverse it, and interrupt it. Those are capability requirements on the human, not features on the software.

the Meaningful Review Test

Five conditions a checkpoint must satisfy to count as review rather than a signature.

01 **competence**

The reviewer can state, without opening the system, what would make this recommendation wrong. If they can only describe what the system considers, they are validating provenance, not exercising judgement. Remediate with case-based training on wrong recommendations, not tool training.

02 **capacity**

The reviewer has time to think about each item. Our working threshold: minutes per item multiplied by items per week should not exceed 40% of the reviewer's working week, since review is rarely their only job. Above that, review degrades to pattern-matching on the interface rather than on the case.

03 **contrary evidence**

The interface shows at least one reason not to accept, alongside the supporting case. A recommendation shown only with its justification is an argument. Shown with its strongest counter-argument, its confidence and the data it could not see, it is a decision aid. Most tools do not meet this by default.

04 **consequence**

The reviewer is measured on decision quality, not throughput. If their manager sees a queue-clearance number and never a decision-quality number, the incentive wins. Remove queue length from reviewer dashboards and put reversal and dispute rates on them instead.

05 **counting**

Override rate is measured per process and per reviewer, reported internally, and investigated at both ends. It is the only cheap instrument that detects a decorative checkpoint, and almost nobody calculates it.

THRESHOLDS FOR THE MEANINGFUL REVIEW TEST, AND HOW TO TEST EACH CONDITION

condition	threshold	how to test it	what a failure looks like
competence	3 of 3 sampled reviewers name a disqualifying condition unprompted	cold interview, no system access, five minutes each	reviewers describe model inputs instead of decision failure modes
capacity	minutes per item times weekly volume stays under 40% of the working week	take the queue log for four weeks and divide	median review time under 30 seconds on decisions with legal effect
contrary evidence	every decide-rung screen shows confidence, one counter-indication and the data not available	screenshot review of each production interface	one recommendation, one justification, one accept button
consequence	reviewer objectives contain a decision-quality measure and no throughput measure	read the actual objectives, not the policy	queue length is visible on the reviewer's screen and quality is not
counting	override rate calculated monthly, per process and per reviewer	query the decision log	nobody can produce the number, so the checkpoint has never been evidenced

On capacity, here is the arithmetic. Suppose an agent proposes 400 actions a week and you want each to get a genuine four minutes. That is 1,600 minutes, or 26.7 hours, roughly two thirds of a working week. Under our 40% threshold, one reviewer supports about 240 items a week at four minutes each. Past that, add a reviewer, narrow what reaches the checkpoint, or move the task to propose. These are our design heuristics, offered because most organisations set no threshold at all.

THE TWO-SIDED OVERRIDE RULE

An override rate below 2% across 100 or more decisions signals a decorative checkpoint. Above 30% signals a model unfit for the rung it occupies, which should be moved down one. Both ends trigger an investigation, not a celebration. These bands are our heuristic; calibrate against your own baseline after two quarters, and set them per process rather than globally.

designing the interface against automation bias

- No pre-focused accept button and no keyboard shortcut that accepts. Symmetry between accept and reject is a decision you have to make on purpose.
- Require a reason code on *accept* as well as reject. It costs three seconds and is the most effective anti-rubber-stamping measure available, because it forces the reviewer to name why.
- Show confidence as a number, and show what the model could not see. Absent data is more decision-relevant than present data, and interfaces almost never display it.

- Insert deliberate friction on the highest-scoring cases. A mandatory second reviewer above a reach score of 4 is cheap at low volumes and prevents the expensive cases.
- Rotate a small share of decisions into a blind review where a second reviewer sees the case without the recommendation, then compare. This is the only reliable way to measure whether the recommendation is anchoring your reviewers.
- Seed the queue with known-bad recommendations at a low rate and record who catches them. Announce the practice in advance so it reads as calibration rather than entrapment.

The blind review and the seeded case are the two practices most likely to be resisted and most likely to tell you the truth. Both come from the mitigation literature, where training focused on automation bias reduced commission errors more reliably than omission errors. Knowing that people miss what the system does not flag is more useful than hoping they will look.

> **A checkpoint you have never measured is a claim. Override rate turns it into a control, and it costs one query against a log you should already be writing.**

Write the reviewer's obligation into their role description in one sentence: you are accountable for this decision whether or not you accepted the recommendation. Then make that true in practice, which means the record has to name them.

CHAPTER 06

accountability, audit trails and the record you keep

The record is the only part of an agentic decision that outlives the model. Design it before you deploy, because you cannot backfill it.

A model gets retrained. A prompt gets edited. A vendor gets replaced. A configuration gets changed by someone who has since left. Eighteen months later an employee contests a decision, and the only surviving thing is the record you wrote at the time. The decision record is not an artefact of governance. It is the product of governance, and everything else is scaffolding around it.

Regulators have started testing whether these records exist. New York City has required annual independent bias audits of automated employment decision tools since 2023. In December 2025 the New York State Comptroller published an audit of how the city's Department of Consumer and Worker Protection enforces that law. Across 32 companies reviewed, the department had identified one instance of non-compliance. The auditors identified at least 17 potential instances in the same set. Separately, 75% of test calls to the city's 311 line about these tools were misrouted and never reached the department.

17 against 1

potential non-compliance instances found by state auditors versus by the enforcing agency, across the same 32 companies

New York State Comptroller, 2025

Read that as a forecast rather than a story about one city. Enforcement of AI-in-employment rules is currently weak and visibly embarrassing to the enforcer. That combination reliably produces more enforcement, not less. The organisations that will be comfortable are those whose records were written for an adversarial reader from the start.

RECORD AND OVERSIGHT OBLIGATIONS THAT ALREADY BIND, WITH DATES

jurisdiction or standard	instrument	effective date	what it requires of your record
European Union	Regulation (EU) 2024/1689, the AI Act; Annex III covers employment and worker management	Annex III standalone high-risk obligations moved to 2 December 2027 by Regulation (EU) 2026/1744, published 24 July 2026	logging, technical documentation, and oversight able to interpret, disregard and interrupt the system
European Union	GDPR Article 22	25 May 2018	a basis for any solely automated decision with legal or similar effect, plus a real route to human intervention and contest
California	automated decision system rules under the Fair Employment and Housing Act	1 October 2025	four years of retention for automated decision system data; applies at five or more employees
Illinois	House Bill 3773, amending the Illinois Human Rights Act	1 January 2026	notice when AI is used in covered employment decisions; no ZIP code proxies
New York City	Local Law 144 on automated employment decision tools	enforcement from 2023	annual independent bias audit, published summary of results, candidate notice
India	Digital Personal Data Protection Act 2023 and DPDP Rules 2025	Act 11 August 2023; Rules notified 13 November 2025, phased	purpose limitation, retention limits and deletion for employee data, with cross-border transfer conditions
Singapore	PDPA 2012; PDPC advisory guidelines on personal data in AI recommendation and decision systems	guidelines issued March 2024	a documented basis for using personal data in decision systems, with heavier oversight where employment access is affected
Saudi Arabia	Personal Data Protection Law and Implementing Regulations; SDAIA AI ethics principles	PDPL in force 14 September 2023, full compliance from 14 September 2024	accountability and transparency duties, fines to SAR 5 million, and a four-tier AI risk classification to map against

jurisdiction or standard	instrument	effective date	what it requires of your record
International	ISO/IEC 42001:2023, AI management systems	published December 2023	an auditable management system: risk assessment, impact assessment, lifecycle control, supplier oversight
United States	NIST AI Risk Management Framework 1.0	released 26 January 2023	a voluntary structure across four functions: govern, map, measure, manage

The pattern across those rows is narrow and consistent. Every instrument wants to know who decided, on what basis, with what notice to the affected person, and how long you kept the proof. One record structure satisfies all of them, which is fortunate, because maintaining separate records per jurisdiction is how multinational HR functions end up with none.

the Eight-Field Decision Record

One row per decision, written at the moment of decision, enough for an external reviewer to reconstruct events without interviewing anyone.

01 **1. identity**

A decision identifier, the process it belongs to, the affected person or population, and the authority statement version it was made under. Without the last item you cannot prove the decision was inside its limits.

02 **2. inputs**

Which data fields were used, from which system, as of which timestamp, including what was unavailable or stale. Absent inputs are the first thing a competent challenger asks about, and reconstructing them a year later is impossible.

03 **3. version**

The model, ruleset or prompt version, with a hash or immutable tag. If you cannot pin the version you cannot rerun the case, and every later explanation is a reconstruction rather than a record.

04 **4. output**

The recommendation exactly as presented to the human, with confidence, ranked alternatives and counter-indications shown. Store what the reviewer saw, not what the system computed. Those differ more often than teams expect.

05 **5. actor**

The named individual who decided, their role, and the delegated authority they acted under. Not a team mailbox, not a service account, not a role title alone. This field is the whole point of the record.

06 **6. disposition**

Accepted, modified or rejected, with a reason code from a controlled list plus free text, and the elapsed time between presentation and decision. Elapsed time reveals a decorative checkpoint and costs nothing to capture.

07 **7. notice**

What the affected person was told, in what language, through which channel, on what date, and what route to contest was offered. Several jurisdictions now require this, and it is the field most often missing.

08 **8. retention**

The retention basis, the review date and the scheduled deletion date. The record has to carry its own clock, because retention duties in different jurisdictions point in opposite directions.

WHERE RETENTION RULES PULL AGAINST EACH OTHER

California requires four years of retention for automated decision system data. India's DPDP framework requires deletion once the purpose is served, with conditions on cross-border transfer. A global employer holds both duties over the same population. Resolve it per record, with a stated basis and a deletion date decided by counsel in advance. Resolving it per incident, under pressure, produces neither compliance nor a defence.

the quarterly record audit, run by someone who did not build the system

- [ ] Sample at least 100 decisions or 5% of the quarter's volume, whichever is larger. For a 2,000-decision quarter that is 100, since 5% of 2,000 is 100 and the floor and the percentage coincide.
- [ ] Check all eight fields are populated. Report completeness by field, not an overall pass rate, because one systematically empty field is the finding.
- [ ] Flag every decision under 30 seconds of elapsed time on a decide-rung process, and investigate the reviewers rather than the cases.
- [ ] Calculate override rate per process and per reviewer, against the two-sided band from chapter five.
- [ ] Verify each referenced model version still resolves to a retrievable artefact. Broken version references are common and fatal to any later defence.
- [ ] Attempt full reconstruction of the three most contestable decisions in the sample from stored artefacts only, and time it.
- [ ] Confirm notice was issued where required, and click the contest route named in it to check it still works.
- [ ] Record the findings, the owner and the remediation date in the same system, so the audit leaves a record of its own.

Vendor arrangements need the same discipline. In Mobley v. Workday, in the Northern District of California, the court accepted that an AI vendor could face direct liability for employment discrimination as an agent of the employer, and in May 2025 granted preliminary certification of a nationwide collective under the Age Discrimination in Employment Act. Whatever the outcome, the theory is live. It does not reduce the employer's exposure. It adds a second defendant with different incentives and different lawyers.

FOUR CLAUSES TO ADD AT RENEWAL

First, a right to the decision log in a machine-readable export you control, retained past termination. Second, notice of model changes that materially affect outputs, before deployment rather than after. Third, a contractual right to conduct or commission bias testing, with the results owned by you. Fourth, cooperation and evidence preservation obligations that survive the contract and bind in litigation. Ask at renewal, when you have negotiating room.

Anchor the structure to an external standard so it survives leadership change. ISO/IEC 42001, published December 2023, gives an auditable AI management system covering risk assessment, impact assessment, lifecycle control and supplier oversight. The NIST AI Risk Management Framework, released 26 January 2023, gives four functions to organise the work: govern, map, measure and manage. Neither substitutes for the record. Both make it defensible, by showing it was produced by a system rather than by a diligent individual.

> **Everything else about an agentic decision is perishable. The model changes, the vendor changes, the configuration changes. The record is what you will be judged on.**

CHAPTER 07

a 90-day pilot and the measures that prove it

Three thirty-day blocks, four numeric gates, and a short list of measures that decorate rather than prove. Start on Monday.

Ninety days is long enough to see behaviour and short enough to stop. The structure below is deliberately conservative in what it lets the system do and deliberately aggressive in what it makes you measure. Most pilots invert that, which is one reason MIT's Project NANDA found roughly 95% of generative AI pilots producing no measurable profit and loss impact. A pilot that cannot fail teaches you nothing you can use.

the 30/30/30 Pilot

Three sequential thirty-day blocks, each with a different permission level and a different question.

01 **days 1 to 30: shadow**

The agent runs on live data and executes nothing, recording what it would have recommended. Humans work as before and do not see the output. Compare recommendations against what humans did. The question is agreement, and the finding that matters most is whether disagreement clusters in any employee group. Uniform disagreement is a tuning problem. Clustered disagreement is a stop.

02 **days 31 to 60: checkpoint**

The agent proposes; a named human accepts, modifies or rejects, with a reason code on every disposition including accept. Nothing executes without the human act. Measure override rate, elapsed decision time, reviewer hours consumed, and the defect rate found in a blind second review of a 10% sample. The question is whether the checkpoint is real.

03 **days 61 to 90: bounded execution**

Release only tasks scoring 2 to 4 on the Reversibility and Reach Test to execute, inside written limits, with breach alerts to a named owner. Keep everything else on the checkpoint. Measure escaped defects, reversal rate, employee-raised disputes and time to detect a deliberately seeded fault. The question is whether your controls catch things without a human watching each one.

04 **day 90: the decision meeting**

One hour, four gates on one page, with the agent owner, the checkpoint owner and someone from employee relations present. Three outcomes only: scale within the current rung, hold and remediate a named gate, or stop. Extending the pilot is not an outcome. It is how programmes avoid a verdict for eighteen months.

THE FOUR-GATE SCORECARD: PASS ALL FOUR TO SCALE, FAIL ONE TO HOLD, FAIL TWO TO STOP

gate	what it measures	threshold	how to compute it
gate 1: agreement	does the agent match competent human judgement	at least 80% agreement on the shadow set, and no employee group whose disagreement rate exceeds the overall rate by more than 10 percentage points	compare shadow recommendations against actual human decisions, then split by group
gate 2: override	is the checkpoint doing work	override rate between 5% and 25% across at least 100 decisions, with median elapsed decision time above 30 seconds	query the decision log by process and by reviewer
gate 3: record	can you reconstruct a contested decision	all eight fields present in a sample of at least 100 records, and three test reconstructions each completed within one working day	run the quarterly record audit against the pilot period
gate 4: burden	did the process improve without hidden cost	cycle time down at least 20% against baseline, with reviewer hours per decision inside a cap set before day 1	measured cycle time plus reviewer time logs; state the cap in the pilot charter

Those thresholds are our design heuristics, not findings. We publish the numbers so you can argue with them rather than invent your own from nothing. Two are worth defending. The 80% agreement floor sits below perfect deliberately, because a system agreeing with humans every time adds no information. The 5% lower bound on override comes from the automation bias literature, where the measured rate of clinicians switching from correct to incorrect after advice was 5.2% at 70% system accuracy. If your humans never change anything, the checkpoint is not absorbing that risk.

choosing the process, and the baseline to capture first

- High volume, so you reach statistical usefulness inside 30 days. Below roughly 200 decisions a month you finish with anecdotes.
- A total of 2 to 5 on the Reversibility and Reach Test, so bounded execution in the third block is defensible.
- An existing measured baseline, or one you can capture in two weeks. A pilot without a baseline produces testimonials.
- A single accountable owner who is not the programme sponsor, so the assessment is not written by the person who funded it.
- Not a hiring, promotion, pay, discipline or termination process. Those are capped at propose and are the wrong place to learn.

1. Cycle time: median and 90th percentile over the previous 90 days, not an average and not a recollection.

2. Volume: decisions per week, with the seasonal pattern noted, so a quiet month is not read as an improvement.
3. Error and rework rate: how often the current process produces a decision later reversed, reopened or disputed.
4. Human effort: hours per decision, measured for two weeks rather than estimated.
5. Distribution: current outcomes split by the employee groups you are legally required to care about. Without this, gate 1 cannot be computed.

measures that prove, and measures that decorate

WHAT TO PUT ON THE DAY 90 PAGE, AND WHAT TO KEEP OFF IT

proves something	decorates	why the second one fails
decisions reversed within 30 days	licences activated	activation measures procurement, not behaviour
override rate and elapsed decision time	prompts or queries per user	interaction volume says nothing about decision quality
cycle time at the 90th percentile	average cycle time	averages hide the tail, and the tail is where complaints originate
employee-raised disputes per 1,000 decisions	user satisfaction with the tool	the tool's users are not the people the decisions are about
reviewer hours per decision	hours saved	hours saved with no redeployment plan is a number nobody can bank
record completeness by field	governance policies published	a policy is an intention; a populated field is evidence

56% of HR functions that do not formally measure the success of their AI investments, with only 16% using ROI measures

SHRM, 2026

The hours saved line deserves particular scrutiny, being the most quoted and least bankable number in this field. Gartner's July 2025 survey of 114 HR leaders found only 7% of organisations giving employees any guidance on how to use time freed by AI. Time saved without a redeployment decision is not a saving. It is dispersal, and it will not appear in any financial statement, which is one reason the value question keeps being asked again.

the day 90 agenda, one hour, one page

- [ ] Read the four gates aloud with the numbers attached. No narrative before the numbers, because the narrative will otherwise explain them away.
- [ ] For any failed gate, name the remediation, the owner and the date. If no owner is available, the gate has failed permanently and the answer is hold.
- [ ] Review the clustered-disagreement analysis from block one with employee relations present. This item is the most likely to be skipped and the most likely to matter.
- [ ] Confirm the agent inventory entry exists, with a named owner, a written authority statement and a stop tested within the last 30 days.
- [ ] Ask the checkpoint owner directly whether human review on this process is real. Their answer is a governance record, so write it down.
- [ ] Decide: scale within the current rung, hold and remediate, or stop. Write the decision, the reasoning and the date into the record system your decisions live in.

Sequence the year from there. One process through 90 days, then two more in parallel on the same structure, then the authority register from chapter three extended across your top twenty tasks. That is a realistic first-year scope for a function at stage two, and it leaves you at a defensible stage three or four where it matters, with the evidence to prove it.

> **A pilot that cannot fail is a purchase with a longer approval cycle. Write the gates before the kick-off, and let the numbers decide.**

A note on method. This paper was assembled by the research team at entomo from public sources, each named and linked so you can check the reasoning rather than accept it. The five frameworks are ours. The thresholds inside them are design heuristics, flagged as such wherever they appear, and meant to be calibrated against your own data after two quarters. The regulatory dates were correct in September 2026, and at least three will move, which is itself an argument for building the record rather than chasing the rule.

REFERENCES

sources

every figure in this document traces to one of the following. links were live at the date of publication.

1 Gartner, 2025 (task-specific AI agents in enterprise apps)
https://www.gartner.com/en/newsroom/press-releases/2025-08-26-gartner-predicts-40-percent-of-enterprise-apps-will-feature-task-specific-ai-agents-by-2026-up-from-less-than-5-percent-in-2025

2 Gartner, 2025 (agentic AI project cancellations)
https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027

3 Gartner, 2025 (fully autonomous AI agents, survey of 360 IT application leaders)
https://www.gartner.com/en/newsroom/press-releases/2025-09-30-gartner-survey-finds-just-15-percent-of-it-application-leaders-are-considering-piloting-or-deploying-fully-autonomous-ai-agents

4 Gartner, 2025 (88% of HR leaders report no significant business value from AI)
https://www.gartner.com/en/newsroom/press-releases/2025-10-28-gartner-survey-shows-88-percent-of-hr-leaders-say-their-organizations-have-not-realized-significant-business-value-from-ai-tools

5 Gartner, 2025 (8% of HR leaders say managers can use AI effectively)
https://www.gartner.com/en/newsroom/press-releases/2025-10-08-gartner-research-finds-only-8-percent-of-hr-leaders-believe-their-managers-have-the-skills-to-effectively-use-ai

6 Gartner, 2025 (the future of AI in HR and the HR operating model)
https://www.gartner.com/en/articles/hr-operating-model

7 SHRM, 2026 (The State of AI in HR, survey of more than 1,900 HR professionals)
https://www.shrm.org/executive-network/insights/state-of-ai-hr-2026-5-critical-insights-chros

8 SHRM, 2026 (Recruiting Executives Benchmarking, more than 4,600 organisations)
https://www.shrm.org/topics-tools/research/recruiting-benchmarking

9 SHRM, 2025 (CHRO Benchmarking: HR-to-employee ratio and HR spend)
https://www.shrm.org/executive-network/insights/more-support-for-hr--talent--chro-benchmarking-survey-insights-

10 Deloitte, 2026 (Global Human Capital Trends, with Oxford Economics)
https://www.deloitte.com/us/en/about/press-room/deloitte-report-winning-organizations-will-build-the-human-advantage.html

11 Deloitte, 2026 (State of AI in the Enterprise, India findings)
https://www.deloitte.com/in/en/about/press-room/indian-enterprises-lead-global-peers-in-at-scale-ai-adoption-across-most-functions.html

12 Mercer, 2026 (Global Talent Trends, nearly 12,000 respondents)
https://www.mercer.com/en-us/about/newsroom/mercer-s-global-talent-trends-2026-report/

13 MIT Project NANDA via Fortune, 2025 (The GenAI Divide: State of AI in Business 2025)
https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/

14 World Economic Forum, 2025 (Future of Jobs Report 2025)
https://reports.weforum.org/docs/WEF_Future_of_Jobs_Report_2025.pdf

15 PwC, 2025 (Middle East Workforce Hopes and Fears Survey, 1,286 employees)
https://www.pwc.com/m1/en/issues/upskilling/hopes-and-fears-2025.html

16 EY survey reported by HR Dive, 2022 (payroll corrections and error costs, 508 respondents)
https://www.hrdive.com/news/employers-make-15-corrections-per-pay-period-on-average-costing-thousands/640276/

17 European Union, 2024 (Regulation (EU) 2024/1689, the AI Act)
https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng

18 European Union, 2026 (Regulation (EU) 2026/1744, the Digital Omnibus on AI)
https://eur-lex.europa.eu/eli/reg/2026/1744/oj/eng

19 EU Artificial Intelligence Act, 2024 (Article 14, human oversight and automation bias)
https://artificialintelligenceact.eu/article/14/

20 Gibson Dunn, 2026 (analysis of the AI Act omnibus and the postponed high-risk deadlines)
https://www.gibsondunn.com/eu-ai-act-omnibus-agreement-postponed-high-risk-deadlines-and-other-key-changes/

21 GDPR, 2016 (Article 22, automated individual decision-making)
https://gdpr-info.eu/art-22-gdpr/

22 New York State Comptroller, 2025 (audit of Local Law 144 enforcement)
https://www.osc.ny.gov/state-agencies/audits/2025/12/02/enforcement-local-law-144-automated-employment-decision-tools

23 New York City Rules, 2023 (Automated Employment Decision Tools, Local Law 144)
https://rules.cityofnewyork.us/rule/automated-employment-decision-tools-2/

24 Jackson Lewis, 2025 (California automated decision system regulations effective 1 October 2025)
https://www.jacksonlewis.com/insights/californias-new-ai-regulations-take-effect-oct-1-heres-your-compliance-checklist

25 National Law Review, 2025 (Illinois HB 3773, effective 1 January 2026)
https://natlawreview.com/article/illinois-anti-discrimination-law-address-ai-goes-effect-1-january-2026

26 Holland and Knight, 2025 (Mobley v. Workday, ADEA collective action certification)
https://www.hklaw.com/en/insights/publications/2025/05/federal-court-allows-collective-action-lawsuit-over-alleged

27 Press Information Bureau, Government of India, 2025 (DPDP Rules 2025 notified)
https://static.pib.gov.in/WriteReadData/specificdocs/documents/2025/nov/doc20251117695301.pdf

28 PDPC Singapore, 2024 (Advisory Guidelines on Use of Personal Data in AI Recommendation and Decision Systems)
https://www.pdpc.gov.sg/-/media/files/pdpc/pdf-files/advisory-guidelines/advisory-guidelines-on-the-use-of-personal-data-in-ai-recommendation-and-decision-systems.pdf

29 SDAIA, Saudi Arabia, 2023 (Personal Data Protection Law and AI ethics regulations)
https://sdaia.gov.sa/en/SDAIA/about/Pages/RegulationsAndPolicies.aspx

30 ISO, 2023 (ISO/IEC 42001:2023, AI management systems)
https://www.iso.org/standard/42001

31 NIST, 2023 (AI Risk Management Framework 1.0)
https://www.nist.gov/itl/ai-risk-management-framework

32 Goddard, Roudsari and Wyatt, JAMIA, 2012 (automation bias systematic review)
https://pubmed.ncbi.nlm.nih.gov/21685142/

33 Goddard, City University of London, 2014 (automation bias and prescribing decision support)
https://openaccess.city.ac.uk/id/eprint/3005/

34 Lyell and colleagues, BMC Medical Informatics and Decision Making, 2017 (automation bias in electronic prescribing)
https://pmc.ncbi.nlm.nih.gov/articles/PMC5356416/

35 Sele and Chugunova, PLOS ONE, 2024 (putting a human in the loop)
https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0298037

about entomo

entomo is an agentic workforce intelligence platform for performance, skills, and talent decisions, used by more than 30 million people worldwide. its agents reason on axon, a living model of work that connects what every role requires with what every person brings, benchmarked against the market and refreshed every month.

the practice behind this document is the same one behind the platform: HR process depth first, intelligence second.

see it against your own hardest role

entomo.co