

designing an internal talent marketplace people use

the operating model, manager incentives and metrics that separate a working marketplace from an empty job board

7 chapters · 15 sources · entomo research

CONTENTS

what is in this ebook

- 01 why marketplaces fill with roles and empty of people

- 02 the manager problem and what the evidence says

- 03 the operating model: gigs, projects, roles and mentoring

- 04 incentives: making talent release pay for managers

- 05 liquidity: the three numbers that decide whether it works

- 06 matching, transparency and the trust the platform needs

- 07 a metric set for the first four quarters

the argument in short

Most internal talent marketplaces launch well and die quietly. The platform goes live, several hundred openings are published, a campaign runs, and usage spikes for about six weeks. Then postings begin to outnumber active seekers. Within two quarters the marketplace has become a notice board, and the real moves are happening the way they always did, through a conversation between two managers. The usual diagnosis is adoption and the usual remedy is more communication. Both are wrong, and the evidence on why is now unusually good.

An internal marketplace is a two-sided market, and its scarce side is not opportunities. It is people willing to be seen looking, attached to managers willing to let them go. Posting a role costs almost nothing. Applying for one costs a relationship with the person who writes your review and sets your pay. In a study of a 200,000-person manufacturer, 41 per cent of employees said they were afraid to apply internally because of what would happen if their manager found out. A quarter believed they needed permission to apply, in a firm whose policy said they did not.

The manager side of that is an economics problem, not a communications problem. A manager evaluated on team output loses that output the day a strong performer moves, while the benefit lands with the receiving team and the firm. Three-quarters of managers in the same study reported hoarding actions. When a manager's own move temporarily removed the incentive to hoard, applications from their team rose 78 per cent, and those who then applied were hired at 49.1 per cent against a firm average of 27.6 per cent. The marketplace was never short of talent. It was short of release.

This document sets out an operating model built on that finding, with five frameworks you can use directly. The hoarding exposure test measures suppression using data you already hold. The four doors decides what work goes on the marketplace and in what order. The release ledger moves the cost of releasing talent off the manager who currently absorbs it. The clearing triangle reduces liquidity to three numbers with thresholds. The four-quarter ladder sequences the metrics so the programme is not killed by reporting the right number too early.

We have used primary sources only, and we name what we set aside. The most repeated statistics in this category come from vendor material with no published method, and several conflict with the benchmarking data. One number should calibrate every ambition below. Across more than 4,600 organisations, the median share of non-executive positions filled externally was 97 per cent. That is the base rate this work starts from. It was compiled by the research team at entomo from published research and public benchmarking data.

this document is provided for general information. it is not legal, accounting or professional advice, and it does not create an advisory relationship.

why marketplaces fill with roles and empty of people

A marketplace has two sides, and most internal marketplaces are carefully built for the side that was never the constraint.

The pattern is predictable. A large organisation buys or builds an internal opportunity platform. Talent acquisition loads several hundred live requisitions on day one, because that is the inventory closest to hand. Registrations climb quickly, because registering costs nothing. Six to eight weeks later the curve flattens, and by the end of the second quarter the ratio has inverted. There are more open postings than people looking at them.

The programme then enters a familiar loop. Someone proposes refreshing the communications plan. Someone else proposes gamification. A third proposes making profile completion mandatory. None of these address what has happened, which is that one side of the market has quietly withdrawn.

the side you did not build for

An internal talent marketplace is a two-sided market. It needs opportunities on one side and available, visible, releasable people on the other. Almost every implementation treats opportunities as the scarce good. In a company of 1,000 or more the opposite is true. What is scarce is a person willing to be seen looking, attached to a manager willing to let them go.

Look at what each side pays to participate. A manager who posts an opening pays almost nothing. The posting is free, it runs in parallel with the external search, and a failed posting costs a few days. An employee who applies pays something immediate. They pay with the relationship they have with the person who writes their review, sets their pay, approves their leave and decides what work they get next. That asymmetry is the whole problem, and no amount of interface design changes it.

None of this is a shortage of appetite. In a global study of nearly 3,900 respondents, 74 per cent said developing skills was important to their organisation's strategy, while only 34 per cent were satisfied with what their organisation invested in them. Nearly half said they were prepared to leave if offered a buyout. A separate survey of 3,375 employees found one in three believed they could have more impact in a different role inside their own organisation. The willingness is measurable. What is missing is a route that does not cost the person taking it.

41%

of employees said they were afraid to apply for internal positions, fearing repercussions if their manager found out

Haegele, American Economic Review, 2026

25%

of employees believed they needed their manager's permission to apply internally, in a firm whose policy explicitly allowed self-initiated applications

Haegele, American Economic Review, 2026

These are not attitudes a launch campaign shifts. They are accurate readings of an incentive structure. The base rate they produce is worse than most business cases assume. In benchmarking covering more than 4,600 organisations, the median share of non-executive positions filled externally rose from 93 per cent in 2025 to 97 per cent in 2026. The median large organisation is not filling a third of its roles internally. It is filling a handful.

WHAT WE SET ASIDE

Several figures that dominate writing on this topic did not survive checking, and we have left them out. The claim that employees stay 41 per cent longer at companies that hire internally comes from a recruiting platform's own report, with no published sample, definition or method. The claim that lateral movers have a 62 per cent chance of staying after three years against 45 per cent for non-movers traces to the same platform's analysis of its own user profiles on a marketing blog. The figure that 39 per cent of roles were filled internally in 2024 traces to secondary aggregation rather than a published dataset, and sits badly against the benchmarking above. Claims that mature marketplaces cut staffing time by a fixed percentage are vendor content with no primary source. One correction is worth naming because it circulates widely: the finding that internal applications rise 123 per cent when hoarding stops comes from a superseded working-paper version of the study we rely on below. The published figure is 78 per cent, and that is the one we use.

three ways it fails, and they are not the same failure

Programmes that stall usually stall in one of three distinct ways. The remedies are different, so the diagnosis matters more than the symptom, which in all three cases looks like low usage.

THREE FAILURE MODES AND WHAT EACH ONE ACTUALLY IS

failure mode	what you observe	what is happening	what changes it
the notice board	high posting volume, low application volume, long time to fill	opportunities are published but nothing is matched, and no one is released	release economics, then matching
the shadow market	healthy platform metrics, but senior and interesting roles are never on it	real moves happen by direct manager arrangement, and the marketplace carries the residual	a posting mandate with a tracked exception rate
the one-way valve	a strong first cohort of applicants, then a sharp drop in repeat applications	people applied, were rejected without process or feedback, and will not return	rejection design and interview coverage

Most organisations have all three at once, in different business units. The notice board is a supply problem. The shadow market is a governance problem. The one-way valve is a trust problem. Only the first is helped by better technology, and even then only after the release economics are fixed.



An internal marketplace does not fail for lack of opportunities. It fails because the cost of looking is paid entirely by the person looking.

the manager problem and what the evidence says

Managers who hoard talent are responding rationally to how they are measured, and there is now direct empirical evidence of how much it costs.

Start with the arithmetic a line manager faces. They are evaluated on the output of their team. A strong performer who moves takes their output with them on the day they leave, and any replacement arrives months later and produces less while ramping up. The benefit accrues to the receiving team and to the firm. This is a textbook externality: the manager bears a concentrated, immediate, certain cost, and the organisation receives a diffuse, delayed, uncertain benefit. Any system that asks managers to absorb that trade voluntarily and repeatedly, without compensation, is asking them to act against their own measured interest. Some will. Most will not, and they are not behaving badly. They are behaving predictably.

what the evidence now shows

Until recently this was an assertion supported by opinion surveys. It is now supported by a study in the American Economic Review combining personnel records, the full history of internal applications and hiring decisions, and original surveys of managers and employees inside a manufacturer of more than 200,000 people. The analysis covers the firm's largest internal labour market, over 30,000 white-collar and management employees, from 2015 to 2019. The employee survey drew more than 15,000 responses at a 50 per cent response rate, and the manager survey reached 62 per cent of managers.

75%

of managers reported taking talent hoarding actions, in a survey covering 62 per cent of the firm's managers

Haegel, American Economic Review, 2026

The manager responses describe the mechanism precisely, and each number points at a different design lever.

- 55 per cent of managers explicitly acknowledged that supporting an employee's development creates a conflict of interest, because developed people are more likely to leave the team.
- 96 per cent believed they had a large impact on their people's careers, but only 36 per cent believed the firm valued that impact.
- 68 per cent said they would be more likely to support development if lost talent were easier to replace.
- Employees reported the hoarding methods directly: suppressed public ratings, restricted access to high-visibility projects and training, and explicit discouragement from applying.

The third of those numbers is the most useful sentence in the literature for anyone designing a programme. Two-thirds of managers named replacement, not values, as the thing that would change their behaviour. That is a supply chain answer to what is usually treated as a culture question. It tells you the intervention is a backfill guarantee, not a workshop.

the natural experiment

The central finding uses a clean identification strategy. When a manager learns they will move to a different team, they briefly stop having a reason to hoard, because the people they would have held on to will no longer be theirs. Their team is, for a short window, exposed to a manager with no hoarding incentive. Nothing else about those employees changes.

78%

increase in internal job applications from a team when the manager's incentive to hoard is temporarily removed, from a baseline application rate of 2.9 per cent

Haegele, American Economic Review, 2026

The quality finding matters more than the volume finding. People who applied only because the hoarding incentive lifted were hired at 49.1 per cent. The firm-wide average for internal applications was 27.6 per cent. Suppressed applicants were substantially better than average, which is what the theory predicts: hoarding is strongest where a person is hardest to replace, so it is strongest on the best people. Your marketplace is not missing marginal talent. It is missing the specific people you would most want to move.

The study also identifies which managers hoard most, which tells you where to look. Hoarding was more common where managers had more performance-related pay, where teams were smaller, and where the function had low visibility. Each additional team member was associated with a 1.4 percentage point reduction in hoarding, and managers of teams of three or fewer showed an 11.4 percentage point increase in hoarding propensity. A manager of six who loses one loses a sixth of their capacity.

There is an equity consequence worth stating plainly, because it changes who is affected. Women in the firm were 22 per cent more likely to place high value on preserving a good relationship with their manager, and 26 per cent more likely to rely on their manager's guidance when making career decisions. The women deterred from applying were more qualified than the men deterred, on both education and past performance. Among those who applied only once the incentive lifted, the pay gap narrowed by 86 per cent. Hoarding is not neutral in its distribution.

the hoarding exposure test

Four signals that measure suppression in your own organisation using data you already hold, with thresholds for action.

01 **signal one: the rating gap**

Compare the private performance rating a manager submits with the public potential rating visible outside the team, for the same person in the same cycle. Systematic suppression of the public rating is the cleanest observable proxy for hoarding. Flag any manager whose public rating sits below their private rating for more than 20 per cent of their team, two cycles running.

02 **signal two: the silent high-performing team**

Calculate the internal application rate per manager per quarter. Isolate managers whose team performance is at or above median but whose application rate sits in the bottom decile. High performance with no outward movement is the signature. A team of eight strong performers producing no internal applications across four quarters is not a coincidence.

03 **signal three: the rotation window**

You already run the experiment the research uses. Whenever a manager announces an internal move, track applications from their team over the next two quarters against that team's own trailing four-quarter baseline. A rise above 30 per cent tells you the previous level was suppressed rather than natural. Run it across all manager moves and you have a continuous internal estimate.

04 **signal four: the permission question**

Add one question to the engagement survey: do you believe you need your manager's approval before applying for an internal role. Report it by manager population, not by business unit average, and read it against your written policy rather than your intent. Anything above 20 per cent is a structural finding.

05 **scoring and action**

Two or more signals flagged for the same manager in consecutive cycles is a management issue, owned by that manager's leader. Two or more signals flagged across more than a quarter of the managers in a function is not a management issue. It is a design fault in how that function is measured, and it belongs with whoever owns the operating model.

WHY CAMPAIGNS DO NOT FIX THIS

Hoarding operates through hidden actions. A manager does not refuse a transfer in writing. They give a lukewarm public rating, decline to nominate someone for a visible project, mention that the timing is difficult, or note that the receiving team is restructuring. None of this is auditable, which is why exhortation, policy statements and unenforceable mobility commitments have such a poor record. You cannot monitor what you cannot observe, so you have to change the payoff instead.



Two-thirds of managers said they would develop their people more if lost talent were easier to replace. That is a logistics answer to a question we keep treating as a values question.

the operating model: gigs, projects, roles and mentoring

The unit of exchange decides everything downstream, so choose it and sequence it before you choose a platform.

Almost every internal marketplace opens with permanent roles, because permanent roles are the inventory talent acquisition already holds. This is the most expensive possible way to start. A permanent role is the highest-cost transaction on the board. It requires the releasing manager to give up a person completely, it changes a reporting line, it triggers a backfill they may not get, and it is irreversible. You are asking a market with no established trust to clear its hardest trade first.

There is a gradient of release cost, and it is steep. Mentoring costs a releasing manager almost nothing. A short gig costs a few hours of attention. A part-time project costs real capacity for a defined period. A permanent move costs the whole person. Sequence your opening by that gradient, not by what is easiest for HR to load.

the four doors

A decision set that assigns any piece of internal work to one of four exchange types, using three variables: duration, share of a person's time, and whether the reporting line changes.

01 door one: mentoring and advice

Under four hours a month, no deliverable owed to the receiving unit, no reporting change, no backfill. The releasing manager is informed, not consulted. Use this door to build the visibility that hoarding suppresses, because restricting access to visible work is one of the primary hoarding mechanisms the research identifies. A mentoring register run outside the line management chain is the cheapest counter available to you.

02 door two: gigs

Under 40 hours in total, under 20 per cent of a person's time, completed within one quarter, with a named deliverable and an accountable owner who is not the person's line manager. No reporting change and no backfill obligation. The manager is notified and has a defined objection window, typically five working days, with a stated reason required. This is the door that generates matching data and manager habit at near-zero release cost.

03 door three: projects

One to six months, 20 to 50 per cent of a person's time, dotted reporting line to the project owner for the duration. Partial bridge cover is owed to the releasing team. Manager agreement is required, with an escalation path to the next level up if agreement is withheld. Track the withholding rate by manager, because this is where hoarding first becomes visible in the data.

04 door four: roles

Permanent, 100 per cent of time, reporting line changes, full backfill guarantee attaches. Do not open this door until the clearing triangle in chapter five is passing for doors two and three, and until you can actually honour a backfill guarantee. Opening door four without a funded backfill mechanism trains every manager in the organisation that the marketplace takes people and gives nothing back.

05 the opening sequence

Open doors one and two together in quarter one. Open door three in quarter two or three, once you have at least 200 completed gigs and a gig fill rate above 60 per cent. Open door four in quarter four at the earliest. If your executive sponsor insists on opening door four immediately, restrict it to a named set of functions where you can guarantee backfill, and say plainly that this is a pilot rather than a launch.

THE FOUR DOORS, WITH THRESHOLDS AND OBLIGATIONS

door	duration	share of time	reporting line	backfill owed	manager's role
mentoring	ongoing, under 4 hours a month	under 3 per cent	unchanged	none	informed only
gigs	up to one quarter, under 40 hours	under 20 per cent	unchanged	none	notified, 5-day objection window with stated reason
projects	1 to 6 months	20 to 50 per cent	dotted line to project owner	partial bridge cover	agreement required, escalation available
roles	permanent	100 per cent	changes	full guarantee	formal release, backfill approved same day

what the research supports, and what it does not

Be clear about the evidential status of this model. The evidence base for internal job posting is strong and peer reviewed. There is a substantial body of work on internal labour markets, on how firms choose between internal and external hiring, and on what happens to people who move. The evidence base for internal gig and project marketplaces specifically is thin to non-existent in the peer-reviewed literature. What exists concerns internal crowdsourcing for ideas rather than internal staffing for work.

AN HONEST NOTE ON DOORS TWO AND THREE

We searched for peer-reviewed evidence that internal gig or project marketplaces improve retention, mobility or staffing cost, and did not find any. The published claims we could locate were vendor material or trade publishing. The design above is therefore an extension from the verified evidence on internal job posting, talent hoarding and rejection, not a finding in its own right. We think the extension is sound, because the release-cost logic is the same and the release cost is lower. But it is an argument, and you should treat it as one.

which work belongs on the market at all

Not every opening should be posted. Research on how firms combine internal and external hiring, using seven years of personnel data from a large investment bank, found that jobs are more likely to be filled internally where performance in the job varies widely between people, and where there are many junior workers relative to senior ones. The first is an information story: when the difference between a good and bad occupant is large, the firm values what it already knows about internal candidates. The second is a supply story.

The same study found no effect for firm-specific skills, which is worth sitting with. The most common argument for internal hiring, that internal people already know our systems and our ways of working,

did not predict internal filling in the data. If you are building the business case on firm-specific knowledge, you are building it on the one mechanism that study failed to find.

the honest picture on lateral moves

Lateral movement is the engine of most marketplace designs, so the evidence on it deserves stating without decoration. A matched-sample study of eight years of personnel data at a large healthcare company compared 2,282 lateral movers with 3,068 matched non-movers. Lateral movers did have a higher subsequent chance of promotion, but only after three to four years. On pay, the effect was around three per cent overall, close to zero for the first two years, and 16 per cent higher than matched non-movers five years after the move. The benefit was driven entirely by moves across sub-functions into genuinely different work, not by moves between business units.

Two findings from the same study cut against the usual pitch. Lateral movement was not associated with improved subsequent performance, so the mechanism is not better job fit. And lateral movers were people who had been in their job longer and performed less well than those who got promoted directly. Lateral mobility is a slow-acting instrument that works through exposure and range, not a quick fix for fit or performance. Design and measure it accordingly.

19 in 100

workers that employers expect could be upskilled and redeployed elsewhere within their own organisation by 2030, out of 59 in 100 who will need training

World Economic Forum, Future of Jobs Report, 2025

That figure is a useful sizing anchor. Employers surveyed across 55 economies expect that if the global workforce were 100 people, 59 would need training by 2030, 29 could be upskilled within their current roles, 19 could be upskilled and redeployed elsewhere in the same organisation, and 11 would be unlikely to receive what they need. Roughly a fifth of your people are the addressable population for doors three and four. Not the whole workforce, and not a rounding error.

before you open any door

- ☐ A named owner for the marketplace who does not report into talent acquisition, because a function measured on requisitions will bias every decision toward door four.

- ☐ A written rule for which work must be posted, with a threshold, and an exception process that is logged rather than informal.

- ☐ A backfill mechanism that has been tested on at least ten real cases before door four opens.

- ☐ A stated position on manager visibility of applications, published before launch rather than discovered by employees afterwards.

- ☐ A baseline measurement of the four hoarding exposure signals, taken before launch so you can tell later whether anything changed.

- ☐ Agreement from finance on how bridge cover is booked, settled as a standing rule rather than negotiated case by case.

- ☐ A decision on what happens to a person's current objectives when they take a gig, because unchanged objectives plus extra work is a tax on participation.

incentives: making talent release pay for managers

A manager who releases a strong performer carries the whole cost and sees none of the benefit, so the design task is to move the cost.

State the transaction precisely, because most programme designs never do. When a strong performer moves from team A to team B, team A loses output immediately and permanently. Team B gains a productive person without paying a recruitment cost or a ramp cost. The firm gains a better allocation of talent. The manager of team A receives, in most organisations, a note of thanks.

Everything the releasing manager pays is concentrated, immediate and certain. Everything they receive is diffuse, delayed and uncertain. This is a price problem. There is a transaction with a real cost and no mechanism for settling it, so the party who bears the cost declines to transact. That is not a failure of culture. It is the correct response to an unpriced externality, and it will persist for exactly as long as the externality is unpriced.

The research points at the lever with unusual clarity. Sixty-eight per cent of surveyed managers said they would be more likely to support their people's development if lost talent were easier to replace. Not if they understood the strategy better. Not if they cared more. If replacement were easier. That is a logistics answer, and it means your primary intervention is a backfill mechanism, not a message.

why exhortation has such a poor record

Hoarding works through hidden action. A manager rarely blocks a move in writing. They give a lukewarm public potential rating while recording a strong private performance rating. They decline to nominate someone for a cross-functional project. They mention that the timing is difficult given the reorganisation. They observe that the receiving team has a reputation for churn. None of this is auditable, and none of it can be disproved.

Because the behaviour cannot be observed, it cannot be reliably punished. Policies requiring managers to support mobility are therefore unenforceable in the specific cases that matter. This is why the intervention has to change what managers are paid for rather than what they are told. The evidence from a randomised field experiment on internal participation at a large public-sector research organisation, covering 7,455 employees, points the same way: recognition delivered through the management line significantly increased participation, while peer recognition had no significant overall effect. The management line is the channel that moves behaviour.

the release ledger

Five instruments that move the cost of releasing talent off the manager who currently absorbs it, with thresholds for each.

01 one: the backfill guarantee

The requisition for a backfill is approved on the day the move is confirmed, not after the person has left. Set a published service level: requisition open within five working days, and the role filled or bridged with interim capacity within 90 days. This is the single highest-leverage instrument in the set, and it is the one most organisations skip because it requires finance to commit headcount before a vacancy exists. If you cannot hold this service level, do not open door four.

02 two: bridge cover

The receiving unit pays a share of the released person's cost back to the releasing team for a defined transition period. Recommended settings: 50 per cent for one quarter on a permanent move, 25 per cent for projects consuming more than 30 per cent of a person's time, zero for gigs and mentoring. This is an internal transfer, not new money, and it must be a standing rule in the operating budget. If it is negotiated case by case it will be negotiated to zero within two quarters.

03 three: release credits

A manager earns credits for each confirmed release, weighted by the performance band of the person released and by whether they were hired into a higher band. Credits buy priority in the next hiring round and priority access to gig talent from the market. Credits expire after four quarters, so they cannot be hoarded in turn. Zero credit is earned for releasing someone in the lowest performance band.

04 four: scorecard weight

Put talent developed and talent released on the manager's own performance review with a stated weight. Ten to 15 per cent is the working range. Below 10 per cent it is ignored in practice. Above 20 per cent it becomes the dominant term and gets gamed, usually by releasing weak performers and by counting mentoring conversations that did not happen. State the weight publicly and keep it stable for at least four cycles, because a weight that moves every year is treated as noise.

05 five: the hoarding review

This is the only negative term, and it should be the only one. Where two or more hoarding exposure signals flag for the same manager in consecutive cycles, that manager's own promotion case pauses pending a review by their leader. Not a sanction, and not automatic. A pause and a conversation. The purpose is to make the cost of hoarding land on the person doing it, which is currently the one place it never lands.

WHO PAYS WHAT, BY DOOR

door	backfill owed to releasing team	bridge cover	release credits earned	manager consent required
mentoring	none	none	low weight	no
gigs	none	none	low weight	no, objection window only
projects	none, capacity returns	25 per cent if over 30 per cent of time	medium weight	yes, with escalation
roles	full guarantee, requisition same day	50 per cent for one quarter	full weight, banded by performance	yes, with escalation

the game you should expect

Any scheme that rewards releasing people will be gamed by releasing the people you least want to keep. This is predictable and you should design for it on day one rather than discover it in quarter three.

- Weight release credits by the performance band of the person released, and award zero for the lowest band.
- Award a credit multiplier where the person is hired into a higher band than they left, because that is the outcome the firm wants and the manager cannot fake it.
- Void the credit if the person returns to a similar level within two quarters, which catches parking moves dressed up as development.
- Publish release counts by manager alongside the performance distribution of those released, so that a manager releasing only weak performers is visible without any accusation being made.
- Do not count internal moves into a manager's own newly created team, which is how release schemes get turned into empire building.

the number that funds all of this

Bridge cover and guaranteed backfill look expensive until they are compared with the alternative, which is external hiring. A study of six years of personnel data covering 5,260 workers at the United States investment banking arm of a financial services firm found that external hires were initially paid around 18 per cent more than workers promoted into the same jobs. They received significantly lower performance evaluations for their first two years. They had around a 61 per cent higher hazard of involuntary exit and a 21 per cent higher hazard of voluntary exit than workers entering through promotion.

State the offsetting finding too, because it is real and it changes how you use the number. External hires who stayed were promoted faster afterwards, and internally promoted workers would only catch up with them after around seven years. External hiring is not simply worse. It is more expensive

and riskier in the near term, and it brings range the internal pool may not have. The case for funding release is that you are currently paying the external premium by default, without ever comparing the two prices on the same page.

IF YOU DO ONLY ONE THING

Approve the backfill requisition on the day the move is confirmed rather than the day the person leaves. It costs the organisation a period of double-counted headcount on a small number of roles. It removes the largest single reason managers give for blocking a move, and it is the only instrument here that changes a manager's answer before they have to give it.

You cannot monitor a hidden action, so you cannot punish it reliably. The only thing left to change is the payoff.

liquidity: the three numbers that decide whether it works

A marketplace works when both sides expect a match, and that expectation is built from three numbers you can observe from week one.

Liquidity is the probability that someone who shows up gets what they came for, within the time they were willing to wait. It is not a measure of size. A marketplace with 40,000 registered profiles and 600 open postings can be completely illiquid, and most are. What matters is whether a participant's expectation of a match is high enough to justify the cost of participating again.

There is one asymmetry that governs everything else. Liquidity fails first on the side with the lower cost of leaving. A hiring manager whose posting attracts nobody has an obvious alternative: they open an external requisition and carry on. Their cost of abandoning the marketplace is close to zero. An employee whose application fails has no equivalent fallback inside the company. They can stop applying, or they can look outside.

That second option is where illiquidity becomes expensive. A study of 9,355 internal rejections at a single large firm over five years found that rejected internal applicants were nearly twice as likely to leave the organisation as employees who were hired internally or who never applied. An illiquid internal marketplace does not simply fail to deliver mobility. It converts your most ambitious people into leavers, which is the precise opposite of the outcome the programme was funded to produce.

the clearing triangle

Three numbers that together tell you whether an internal marketplace is clearing, with thresholds and the failure each one predicts.

01 number one: fill rate, the demand side

The share of posted opportunities filled within their own stated window. Below 40 per cent, managers stop posting, because posting has become a performance rather than a route to getting work done. The working target is 60 to 75 per cent. Above 90 per cent is also a warning sign, not a success: it usually means managers are only posting work they have already decided who will do, which is the shadow market showing up in your metrics as excellence.

02 number two: match rate, the supply side

The share of active seekers who receive at least one shortlist or interview within 90 days of their first application. This is the number almost nobody tracks and the one that predicts collapse. Below 25 per cent, the supply side concludes the market is not for them and stops returning. Target 40 per cent or better. Measure it per person, not per application, because one person making eight applications and getting nothing is a very different signal from eight people getting one each.

03 number three: time to clear

Median days from posting to filled, tracked separately by door. Gigs should clear in under 14 days, projects in under 30, roles in under 45. These are not efficiency targets. They are relevance thresholds. Work that takes longer than this to staff internally has already been done another way, usually by the manager giving it to whoever was nearest, which teaches everyone that the market is slower than the corridor.

04 the reading

All three passing means the market clears and you can open the next door. Fill rate passing with match rate failing means you are serving managers at the expense of employees, and your supply side will be gone within two quarters. Match rate passing with fill rate failing means you have willing people and unattractive or badly specified work. Time to clear failing on its own means a process problem, usually approvals, and it is the easiest of the three to fix.

CLEARING THRESHOLDS BY DOOR

door	fill rate floor	fill rate target	median time to clear	review trigger
gigs	40 per cent	60 to 75 per cent	under 14 days	two consecutive months below floor
projects	35 per cent	50 to 65 per cent	under 30 days	two consecutive months below floor
roles	30 per cent	40 to 60 per cent	under 45 days	one quarter below floor
all doors, match rate	25 per cent	40 per cent or better	first match within 90 days	any month below floor

the concentration check

Add one number to the triangle as a health check rather than a threshold: the count of distinct managers who posted at least once in the quarter. If more than 30 per cent of your postings come from fewer than 10 per cent of your managers, you do not have a market. You have a cluster of enthusiasts, and their enthusiasm is doing work that the design should be doing. Enthusiast clusters look healthy in aggregate reporting and collapse when one or two people change jobs.

A reasonable trajectory is 10 per cent of people managers posting at least once by the end of quarter two, and 25 per cent by the end of quarter four. Report the distribution, not the total. A single business unit posting 400 gigs while eleven others post nine between them is a fact your executive summary should contain.

solving the cold start

Every two-sided market has a cold start problem, and there are only two ways to solve it: guarantee supply or guarantee demand. In an internal marketplace, guarantee demand. You control it directly, because the organisation already owns the work. You cannot guarantee supply, because supply depends on managers releasing people and employees feeling safe enough to look.

- Set a posting mandate with a numeric threshold. Any project requiring more than 20 person-days of effort must be posted for five working days before it can be staffed by direct approach.
- Require a logged exception, approved one level up, for any internal appointment made without a posting. Track the exception rate as a headline metric rather than an operational one.
- Seed the first quarter deliberately. Ask each function's leader for five real gigs with named deliverables before launch. Nothing kills a marketplace faster than an opening inventory of vague development opportunities that lead nowhere.
- Do not fill the board with external requisitions to create the appearance of volume. It inflates postings, depresses fill rate, and teaches employees that the marketplace is a job board they are competing against the outside world on.

what liquidity is not

Registered users, profile completion rates, logins and page views are inputs. They tell you the communications worked. They are also the numbers most likely to appear in a steering committee pack, because they move early and they move upward. A marketplace can post excellent numbers on all four while clearing almost nothing. If you report those figures at all, report them beside the clearing triangle, never instead of it.

49.1% vs 27.6%

hiring rate for internal applicants whose applications were previously suppressed by manager hoarding, against the firm-wide average hiring rate

Haegel, American Economic Review, 2026

Keep that comparison in view when liquidity looks poor. Suppressed supply is not weak supply. The people your market is not reaching were hired at nearly twice the average rate once the suppression lifted. Illiquidity in an internal marketplace is rarely a shortage of capable people. It is a queue of capable people who have concluded, correctly, that stepping forward is not worth what it costs them.

A manager whose posting fails opens a requisition instead. An employee whose application fails opens a browser.

matching, transparency and the trust the platform needs

People will not put their ambition into a system whose workings they cannot see, and once they stop, better matching does not bring them back.

Trust is not a soft consideration in this design. It is the input to supply. Every decision about who sees what either raises or lowers the cost of being seen looking, and that cost is the thing suppressing your market. Treat visibility rules as economic parameters, because that is what they are.

the visibility question

There is one design decision that matters more than the rest of the platform configuration combined: can a line manager see that their team member has applied. There are three defensible answers, and the trade-offs are genuine.

THREE VISIBILITY SETTINGS AND WHAT EACH ONE COSTS

setting	what the manager sees	effect on supply	effect on release	use when
full visibility	every application at submission	severe suppression, the 41 per cent stay away	manager can plan early, and can also block early	almost never, unless required by law or works council agreement
notify on shortlist	only once the receiving manager shortlists	moderate, the decision to look is private	manager learns in time to plan, too late to suppress quietly	the recommended default
notify on offer	only when an offer is made	strongest supply, lowest trust cost	manager is surprised, which hardens hoarding for the next case	early stage only, or where hoarding exposure is severe

Notify on shortlist is the right default for most organisations. It gives the releasing manager time to plan and preserves their dignity, which matters more than it sounds, while removing the veto that operates at the point of application. Notify on offer buys short-term supply at the cost of long-term manager cooperation. Full visibility is the setting that produces the empty marketplace, and it is often chosen because it feels transparent. It is transparent in one direction only.

posting versus quietly slotting someone in

The competitive route is not merely fairer than the manager-arranged route. Research on internal hiring processes at a large firm, published in *Administrative Science Quarterly*, compared internal

hires who came through a posted competitive process with those a manager identified personally and slotted into an open job. Those who came through the posted process performed better, were less likely to leave the firm, and earned more than those who were slotted.

This has a direct governance implication. If the marketplace carries only the roles nobody has already decided, it will produce worse hires than the corridor, and the comparison will be used against it. The rule that follows is simple to state and hard to hold: any internal appointment made without a posting requires a written exception approved one level above the hiring manager, and the exception rate is reported to the executive team every quarter. Above 30 per cent, you are running a shadow market with a marketplace attached to it for appearances.

the rejection problem is the retention problem

Most internal marketplace designs spend their attention on matching and almost none on rejection. The evidence suggests the ratio should be closer to the reverse. In the study of 9,355 rejections cited earlier, rejected internal applicants were nearly twice as likely to leave. Two findings from the same work point directly at design changes.

First, candidates who were rejected after reaching the interview stage were around half as likely to exit as those rejected earlier in the process. The interview itself carries information: it tells the person they were close, and it gives them something to act on. Second, a rejected candidate's likelihood of leaving was cut in half if they were passed over in favour of another internal candidate rather than an external hire. The message received is not only that you were not chosen. It is what the organisation chose instead of you.

1. Interview more internal candidates than the shortlist strictly requires. The interview is functioning as a retention instrument, and it is cheaper than a replacement hire.
2. Tell rejected internal applicants who was appointed and whether that person came from inside the company. This is uncomfortable and it halves the attrition effect, which is a trade worth making.
3. Give every rejected internal applicant a named next step within ten working days, owned by a person rather than a system. Not encouragement. A specific gig, a specific development action, or a specific role to target and what would need to change.
4. Track attrition among rejected internal applicants as a standing metric. If it exceeds your general voluntary attrition rate by more than 50 per cent, the rejection process is the problem, not the matching engine.

the internal candidate guarantees

- ☐ You will not be asked for your manager's permission to apply, and the platform will not ask for it on your behalf.
- ☐ Your manager will be told when you are shortlisted, not when you apply.
- ☐ You will receive a decision within a stated number of days, and the number will be published before you apply.
- ☐ If you reach interview and are not selected, you will be told what was missing, by a person, within ten working days.
- ☐ You will be told whether the role went to an internal or an external candidate.
- ☐ Applying will not be recorded in your performance file, and no performance rating may reference an internal application.
- ☐ If your manager objects to a gig or project, they must give a reason in writing that you can see, and you may escalate one level.

what the matching engine owes people

Skills data in most large organisations is thin, inconsistent and out of date. Gartner research found that only eight per cent of organisations have reliable data on the skills their workforce currently possesses and which of those skills matter most to business results. This has a consequence for interface design that is routinely ignored: a matching engine built on unreliable data must not present its results with false precision. Explainability here is not an ethical add-on, it is a supply mechanism. People return to a system whose logic they can work with, and abandon one that appears to judge them by rules they cannot see. Given how weak the underlying data usually is, the honest interface is also the more effective one.

THE PRECISION TRAP

A match score of 87 per cent, generated from a skills inventory that is two years old and half self-reported, is a confident-sounding number about nothing. It is also unfalsifiable to the person reading it, which makes it feel arbitrary the first time it is wrong. Show the two or three signals that drove the match instead, and show what would change the answer. A person who can see that a match rests on three years of adjacent project experience and one certification can act on that. A person shown 87 per cent can only accept or resent it.

a metric set for the first four quarters

Most internal marketplaces are killed by reporting the right number in the wrong quarter, so sequence the measures to match what can be true yet.

Here is how these programmes usually die. In the second steering committee, the finance sponsor asks what has been saved. The team, wanting to be helpful, produces a cost-avoidance figure based on four internal fills. It is a small number. It is also an honest number, and it is fatal, because it establishes cost avoidance as the measure of the programme at the exact point in its life when cost avoidance cannot possibly be large. Nine months later the budget is reallocated.

The remedy is to agree, before launch, what will be reported in each of the first four quarters, and to hold that agreement in writing. Sequence matters more than selection here. Most of these metrics are reasonable. Reporting them in the wrong order is what does the damage.

the four-quarter ladder

A reporting sequence that matches each quarter's headline metric to what the programme can plausibly have achieved by then.

01 quarter one: can people see the market

Report supply readiness only. Share of the workforce with a verified skill or experience profile, target 60 per cent in the launch population. Share who have viewed at least one opportunity, target 25 per cent. Distinct managers who have posted at least once, target 10 per cent of people managers. Baseline reading on all four hoarding exposure signals. Report no outcome metrics at all. The honest headline for quarter one is that the market exists and people can find it.

02 quarter two: does the market clear

Report the clearing triangle for doors one and two. Fill rate, match rate, median time to clear, each against the thresholds in chapter five. Add the concentration check. This is the quarter where you find out whether you have a market or a notice board, and it is the right quarter to find out, because doors three and four have not opened yet and the cost of correcting is still low.

03 quarter three: do managers release

Report release behaviour. Number of confirmed releases by door. Backfill service level attainment, measured as the share of role moves where the requisition was approved on the day of confirmation and filled or bridged within 90 days. Manager objection rate on gigs and projects, with the distribution by manager. Hoarding exposure signals against the quarter-one baseline. This is the quarter that predicts whether the programme survives contact with door four.

04 quarter four: are the moves good, and cheaper

Now report economics and quality. Internal fill share of eligible roles. Cost comparison against the external route for the same role families. Attrition among rejected internal applicants against general voluntary attrition. Twelve-month performance distribution of internal movers against matched non-movers. This is also the first quarter in which a cost-avoidance number means anything, because it is the first quarter in which the denominator is real.

05 the standing rule

Do not report a quarter-four metric in quarter one, even when asked directly, and especially when asked directly. The correct response to an early request for cost savings is to show the agreed ladder and say which rung the programme is on. If that answer is not acceptable to the sponsor, you have learned something important about the programme's survival odds while it is still cheap to learn it.

THE LADDER, WITH THRESHOLDS

quarter	headline question	primary metrics	threshold for green
one	can people see it	profile completeness, opportunity views, distinct posting managers	60 per cent, 25 per cent, 10 per cent of people managers
two	does it clear	fill rate, match rate, time to clear, posting concentration	fill above 60 per cent on gigs, match above 25 per cent, gigs under 14 days
three	do managers release	confirmed releases, backfill service level, objection rate, hoarding signals	backfill service level above 80 per cent, objection rate under 20 per cent
four	are the moves good	internal fill share, cost against external route, rejected-applicant attrition, 12-month performance	rejected-applicant attrition within 50 per cent of general attrition

pair every metric with its counter-metric

Each headline number in this set has an obvious way to be made to look good. Publish the pair, not the number, and the incentive to game it mostly disappears because both halves have to move together.

METRICS AND THE COUNTER-METRIC THAT KEEPS THEM HONEST

metric	how it gets gamed	counter-metric
number of postings	posting work already allocated to a named person	fill rate, plus share of postings with more than one applicant
internal fill rate	filling roles internally with weak matches to hit a target	12-month performance of internal movers against matched non-movers
applications received	campaign-driven applications with no intent behind them	attrition among rejected internal applicants
releases by manager	releasing only weak performers to earn credits	performance band distribution of released people
profile completion	mandating completion in a compliance push	share of profiles updated voluntarily in the last six months

set a target you can actually hit

Ambition here should be calibrated against the base rate, not against the case study. In benchmarking data from more than 4,600 organisations, the median share of non-executive

positions filled externally was 97 per cent. A programme that promises 40 per cent internal fill in year one is promising something the median organisation is roughly thirty percentage points away from. It will miss, and the miss will cost it the mandate.

97%

median share of non-executive positions filled externally, across more than 4,600 organisations, up from 93 per cent the previous year

SHRM Recruiting Executives Benchmarking, 2026

Ten per cent of eligible roles filled internally in year one is a strong result against that base rate. Define eligible narrowly and publish the definition, because an undefined denominator will be redefined by whoever is unhappy with the numerator. And commit to the long horizon in writing, because the evidence says the returns arrive late. In the matched-sample study of lateral movers, the pay premium was close to zero for two years and reached 16 per cent only at five years. A programme measured on twelve-month returns is being measured on the part of the curve where nothing has happened yet.

the failure mode to name in advance

There is a documented pattern worth putting in front of your executive sponsor before you start. When large employers announced skills-based hiring and dropped degree requirements, researchers examined 11,300 roles against the career histories of 65 million workers. Removing a requirement did raise non-degree hiring by around 3.5 percentage points in the roles affected. But only 3.6 per cent of roles actually changed, so the net effect came to 0.14 percentage points, which the authors summarised as not even one in 700 hires. Around 45 per cent of the firms studied had changed their stated policy and not their behaviour at all.

That is the exact failure this document is written against. The announcement is cheap, the platform is buyable, and neither one moves the number unless the economics underneath change. The same research found that the firms which did change behaviour saw real results: workers hired into formerly degree-required roles gained around 25 per cent in salary, and at those firms two-year retention was ten percentage points higher than at firms that announced and did nothing.

the questions to ask at each quarterly review

- ☐ Which rung of the ladder are we on, and are we reporting metrics from that rung only.
- ☐ What is the match rate, measured per person rather than per application, and is it above 25 per cent.
- ☐ How many distinct managers posted this quarter, and is the distribution widening or narrowing.
- ☐ What proportion of internal appointments were made without a posting, and was each exception logged and approved one level up.
- ☐ What is attrition among internal applicants we rejected, compared with general voluntary attrition.
- ☐ Did we honour the backfill service level on every role move, and if not, which ones and why.
- ☐ Which managers have flagged on two or more hoarding exposure signals for two consecutive cycles, and has their leader had the conversation.
- ☐ What did we say we would report next quarter, and is anything about to be reported early under pressure.

The final test is simpler than any of the numbers above. If the platform were switched off tomorrow, would managers still release people, and would employees still put their names forward. If the answer is no, the marketplace was software rather than an operating model, and the metrics were measuring a product's usage rather than an organisation's behaviour. Everything in this document is aimed at making the answer yes, which is why so little of it concerns the platform at all.

REFERENCES

sources

every figure in this document traces to one of the following. links were live at the date of publication.

- 1 Haegele, Talent Hoarding in Organizations, American Economic Review 116(8), 2026
<https://www.aeaweb.org/articles?id=10.1257/aer.20220264>

- 2 Haegele, Talent Hoarding in Organizations, open working paper (version of record for figures cited), 2025
<https://arxiv.org/abs/2206.15098>

- 3 Dlugos and Keller, Turned Down and Taking Off: Rejection and Turnover in Internal Talent Markets, Academy of Management Journal 64(1), 2021
<https://pure.psu.edu/en/publications/turned-down-and-taking-off-rejection-and-turnover-in-internal-tal/>

- 4 Cornell University, Rejected Internal Applicants Twice as Likely to Quit, 2021
<https://www.news.cornell.edu/stories/2021/08/rejected-internal-applicants-twice-likely-quit>

- 5 Keller, Posting and Slotting: How Hiring Processes Shape the Quality of Hire and Compensation in Internal Labor Markets, Administrative Science Quarterly 63(4), 2018
<https://journals.sagepub.com/doi/abs/10.1177/0001839217736045>

- 6 Bidwell, Paying More to Get Less: The Effects of External Hiring versus Internal Mobility, Administrative Science Quarterly 56(3), 2011
https://faculty.wharton.upenn.edu/wp-content/uploads/2012/03/Paying_More_ASQ_edits_FINAL.pdf

- 7 Bidwell and Keller, Within or Without? How Firms Combine Internal and External Labor Markets to Fill Jobs, Academy of Management Journal, 2014
https://faculty.wharton.upenn.edu/wp-content/uploads/2016/11/Within-or-without_final.pdf

- 8 Bidwell and Keller, Stepping Sideways to Step Up: Lateral Mobility and Career Advancement Inside Organizations, Management Science 71(1), 2025
<https://faculty.wharton.upenn.edu/wp-content/uploads/2026/07/Stepping-sideways-to-step-up.pdf>

- 9 Sigelman, Fuller and Martin, Skills-Based Hiring: The Long Road from Pronouncements to Practice, Burning Glass Institute and Harvard Business School, 2024
<https://www.burningglassinstitute.org/research/skills-based-hiring-2024>

- 10 Gartner, HR Research Finds Current Talent Management Efforts Inhibit Employee and Organizational Performance, 2024
<https://www.gartner.com/en/newsroom/press-releases/2024-09-18-gartner-hr-research-finds-orgs-current-talent-management-efforts-inhibit-optimal-employee-and-org-performance>

- 11 SHRM, 2026 Recruiting Executives Benchmarking: Attracting Critical Talent, 2026
<https://www.shrm.org/topics-tools/research/recruiting-benchmarking>

- 12 World Economic Forum, Future of Jobs Report 2025, 2025
<https://www.weforum.org/publications/the-future-of-jobs-report-2025/>

- 13 Schrage, Schwartz, Kiron, Jones and Buckley, Opportunity Marketplaces, MIT Sloan Management Review and Deloitte, 2020

- 14 Gallus, Jung and Lakhani, Recognition Incentives for Internal Crowdsourcing: A Field Experiment at NASA, Harvard Business School, 2020
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3487992
 - 15 LinkedIn, Global Talent Trends 2020, cited here only as a claim we set aside, 2020
<https://www.linkedin.com/business/talent/blog/talent-management/employees-stay-41-percent-longer-at-companies-that-do-this>
-

about entomo

entomo is an agentic workforce intelligence platform for performance, skills, and talent decisions, used by more than 30 million people worldwide. its agents reason on axon, a living model of work that connects what every role requires with what every person brings, benchmarked against the market and refreshed every month.

the practice behind this document is the same one behind the platform: HR process depth first, intelligence second.

see it against your own hardest role

entomo.co

