

EBOOK

the skills-based organisation, practically

a staged first-year guide from role architecture to skills-based staffing, with what to do in each quarter

7 chapters · 12 sources · entomo research

CONTENTS

what is in this ebook

- 01 why skills programmes stall after the taxonomy

- 02 quarter one: role architecture and the work you have

- 03 quarter two: choosing a taxonomy you can maintain

- 04 quarter three: inference, validation and confidence thresholds

- 05 quarter four: skills-based staffing and the first real decisions

- 06 rewiring job architecture, pay bands and promotion

- 07 the metrics that show it is working, and the ones to drop

the argument in short

Almost every large enterprise now has a skills programme. Very few can name a staffing decision that changed because of it. The gap is not a failure of ambition or of technology. It is a sequencing error, and it repeats with enough regularity to be predictable.

The usual sequence puts the taxonomy first. A library gets built, profiles get populated, a dashboard shows coverage climbing, and the programme reports success against a measure that no operating leader asked for. Meanwhile requisitions are still written against job titles, pay is still anchored to grades, and promotion still means a new title. Nothing downstream consumes the skills data, so the data decays and the programme quietly ends.

The correction is not more rigour applied to the same artefact. Organisations that get somewhere do three things differently. They decide which decision will change before they decide what to call anything. They treat the skill model as infrastructure with a maintenance budget rather than a project with a completion date. And they are explicit, in writing, about how much they trust each piece of data and what it is therefore allowed to influence.

This guide runs the sequence in the other direction. It starts from the work rather than the vocabulary, it treats the taxonomy as a maintenance commitment rather than a purchase, it sets explicit confidence thresholds before inferred data touches a decision, and it puts skills into one real staffing decision inside the first year.

Each quarter is written with entry criteria, the work itself, named owners, exit criteria and the failure most likely to occur. Three named frameworks carry the harder judgements: the taxonomy fitness test, the confidence ladder and the staffing switch.

The intended reader is a head of talent, workforce strategy lead or HR technology owner at an enterprise above one thousand employees who has either started a skills programme that has stalled, or is about to start one and would like not to.

Nothing here requires a new platform, a reorganisation, or a change to the job architecture in year one. It requires a decision owner, forty roles described properly, one honestly chosen taxonomy, published confidence thresholds, and the willingness to let one hiring manager see a shortlist they would not otherwise have seen.

this document is provided for general information. it is not legal, accounting or professional advice, and it does not create an advisory relationship.

why skills programmes stall after the taxonomy

the library gets built, the profiles get populated, and not one staffing decision changes.

The pattern is consistent enough to describe in advance. A skills programme is approved. A taxonomy is selected or built. Role profiles are populated, often through a large interview exercise with subject matter experts. A coverage dashboard appears, showing the percentage of the workforce with a completed skill profile. That number climbs, and the programme is reported as on track.

Somewhere in month seven or eight, a quieter fact surfaces. Requisitions are still written against job titles. Pay bands still key off grades. Promotion still means a title change. The skills data is real, it is reasonably accurate, and nothing consumes it.

93% / 20%

of HR executives said moving away from a job-centric model is important, against the 20% who said their organisation was very ready to do it

Deloitte, The skills-based organization, 2022

The distance between those two numbers is the subject of this guide. Intent is close to universal. Readiness is not, and the gap does not close by restating the intent more forcefully.

10%

of HR executives said their organisation effectively classifies skills into a taxonomy, against 85% who had classification efforts underway

Deloitte, The skills-based organization, 2022

That gap between effort and effectiveness is usually read as evidence that the work is unfinished. A more useful reading is that the work was aimed at the wrong target. A taxonomy is a controlled vocabulary. It resolves what things are called. It does not, on its own, change any decision.

This is worth stating plainly because the taxonomy is usually the most visible and most expensive artefact the programme produces. It absorbs the budget, it absorbs the executive attention, and it produces a deliverable that can be demonstrated. None of that makes it the thing that changes an outcome.

the evidence from skills-based hiring

The most rigorous look at whether skills-based intent translates into skills-based practice came from Harvard Business School and the Burning Glass Institute. Examining firms that had publicly committed to skills-based hiring, they found that dropping a degree requirement from a posting raised the share of hires without a bachelor's degree by roughly 3.5 percentage points.

In aggregate that reached fewer than one in seven hundred United States hires in 2023. Of the firms that had made public commitments, 45 per cent changed nothing measurable at all.

READ THIS CAREFULLY

The finding is not that skills-based hiring does not work. It is that removing a requirement from a job posting is not skills-based hiring. The posting changed. The screening, the shortlisting and the interview did not.

three stalls, and how to tell them apart

the stall diagnostic

run this before restarting a programme, because the three stalls need different repairs.

01 the artefact stall

The taxonomy became the deliverable. Symptom: the programme reports coverage, completeness and skill counts, and cannot name a decision that changed. Repair: pick one decision and work backwards from it.

02 the evidence stall

Skill data exists but nobody trusts it, usually because it is entirely self-reported. Symptom: managers override matches and hiring teams ignore internal candidates. Repair: set confidence thresholds and validate a sample.

03 the decision stall

The data is trusted, and the downstream processes were never rewired to consume it. Symptom: requisitions, pay and promotion still key off titles. Repair: change one requisition template and one pay rule, not the whole architecture.

Most stalled programmes are in the artefact stall and believe they are in the evidence stall. They respond by buying better inference, which produces more accurate data that still nothing consumes.

A taxonomy tells you what to call things. It cannot tell you what is true, and it cannot make a decision.

what the stall costs, in the currency a finance director uses

The artefact stall is rarely killed outright. It is defunded quietly, usually in the second budget cycle, when someone asks what changed and the answer is a number about the model rather than about the business. By then three costs have already been incurred.

cost	how it shows up	who absorbs it
sunk build cost	the interview programme, the consulting engagement and the platform configuration behind a model nothing consumes	the HR budget, once
decay cost	profiles built at significant expense that are wrong within eighteen months because no refresh cadence was funded	the HR budget, annually, invisibly
credibility cost	the next workforce initiative is met with a question about what happened to the last one	the function, for several years

The third is the one that matters. A defunded skills programme makes the next attempt harder, because the organisation has learned that this category of work produces dashboards. That is the real argument for putting a decision inside year one rather than year three.

why the job is the wrong container

Deloitte's research on the skills-based organisation surveyed workers and executives across ten countries. It found that 71 per cent of workers already perform work outside their job descriptions, and that only 24 per cent said they do the same work as other people holding the identical job title.

Three in four workers do not do the same work as their same-titled colleagues. A job title is not a category. It is a label sitting on top of a distribution, and any analysis that treats it as a category inherits that error.

THE ANALYTIC CONSEQUENCE

Any capability gap calculated at job-title level carries the variance inside the title. If three in four people with the same title do different work, a gap of twenty per cent for that title tells you almost nothing about any individual, and cannot be used to plan learning, hiring or redeployment for any of them.

That is the case for a smaller unit. The rest of this guide is about building one that can be maintained, trusted and used, in that order, inside four quarters.

quarter one: role architecture and the work you have

start from what people do, not from what the market calls it, and start with the roles that carry your headcount.

QUARTER ONE AT A GLANCE

	what it is
entry criteria	one named executive sponsor, and one decision the programme is expected to change within twelve months
the work	decompose 30 to 50 roles into responsibilities, tasks, KPIs and skills
owners	line leaders own role content, a central team owns the method and the format
exit criteria	every selected role has 8 to 12 skills with a required proficiency, and each skill traces to a task
common failure	attempting the full role catalogue, which stalls somewhere around role forty

Quarter one produces the structure everything later depends on. Get it wrong and the taxonomy choice, the inference model and the staffing change all inherit the error.

pick the roles that carry the weight

Take the 30 to 50 roles that account for most of your population, and add any role that is critical, scarce or currently hard to fill. Do not start from the full job catalogue. A map of nine hundred roles is not a map, it is a second catalogue, and it will not be maintained.

Coverage is not the goal in quarter one. Depth is. A well-built architecture for forty roles supports a real decision. A shallow one for nine hundred supports a dashboard.

who owns what, decided before the first session

Most role architecture work fails on ownership rather than method. If the central team writes role content, line leaders do not recognise their own work in it and will not defend it. If line leaders own the format as well as the content, you get forty incompatible documents.

artefact	owned by	reviewed by
role content: responsibilities, tasks, KPIs	the line leader accountable for the role	the central team, for format only
the method, template and proficiency scale	the central team	nobody; it is not negotiable per role
skill names and their mapping	the central team	the line leader, for recognisability
the decision the programme will change	the executive sponsor	the steering group, quarterly

Write this down before the first role session. Ownership disputes discovered in month five cost a quarter.

the decomposition, and the test each layer must pass

the five-layer decomposition

each layer has one job and one test, and a layer that fails its test is decoration.

01 role

The scope a person is accountable for. Test: you can name the outcome that fails if nobody holds it.

02 responsibility

A durable area of accountability, usually three to seven per role. Test: it survives a change of tooling and a change of team structure.

03 task

The observable unit of work. Test: someone could watch it being done and know it was done.

04 KPI

The evidence a responsibility was discharged. Test: it moves when the work is done well and stays flat when it is not.

05 skill

The capability a task requires, at a stated proficiency. Test: you can describe what a person at level three does that a person at level two cannot.

The chain matters more than any single layer. Skills without tasks float free, which is why most skills inventories cannot be used for anything. Tasks without KPIs cannot be prioritised, because nothing indicates which ones carry the outcome.

The most common error in this step is naming tools where skills belong. A person does not have a skill in a particular applicant tracking system. They have a skill in requisition management, evidenced by working in that system. Tools change on procurement cycles. Capabilities do not, and a model built on tool names has to be rewritten every time a contract is renegotiated.

how to run a role session in ninety minutes

1. Invite the manager, two consistently strong performers, and one person who joined within the last six months. The new joiner still remembers what surprised them.
2. Spend the first twenty minutes listing tasks only. Ask what happened last week, not what the role involves.
3. Group tasks into responsibilities. If you need more than seven groups, two roles are probably wearing one title.
4. Attach one KPI per responsibility. Where there is no measurable evidence, write down why and move on rather than inventing a metric.
5. For each task, name the capability required and the depth required. Depth is the harder and more useful half.
6. Compare the result against the existing job description, and record what appears in the document but never came up in the room.

THE CONSTRAINT THAT DOES THE WORK

Cap each role at 8 to 12 skills. A forty-skill profile describes a person's whole working life. Under that treatment every role looks distant from every other role, and no adjacency calculation will produce a usable answer.

The question that forces the cap is simple. Which skills would make an otherwise capable person fail in this role if they lacked them? Everything else is context, and context belongs in the description rather than the model.

when one title is really two roles

Role sessions frequently surface a title that cannot be described coherently, because two different jobs are wearing it. The signs are consistent.

- the group cannot agree on the top three tasks, and the disagreement splits along team or geography lines
- the responsibility list runs past seven and cannot be consolidated without losing something real
- two people in the room describe the same title with almost no overlap in their week
- the skill list exceeds fifteen and every attempt to cut it removes something that a subset genuinely needs

When this happens, split the role in the model and leave the title alone. You do not need to change anyone's title to describe their work correctly, and attempting to do so in quarter one will attract resistance the programme has not yet earned the standing to absorb.

what to write down about proficiency

Proficiency scales fail when the levels are not behaviourally distinct. If two experienced assessors cannot independently place the same person at the same level, the scale is decorative and every downstream calculation inherits the noise.

Write each level as what a person at that level produces, decides and escalates. Three well-anchored levels beat five vague ones, and three is enough to support a staffing decision.

level	produces	decides	escalates
1 · applies	standard output under supervision, following an established method	which method to apply from a known set	anything non-standard
2 · owns	complete output independently, including non-standard cases	how to adapt the method to the situation	cross-functional trade-offs and precedent-setting calls
3 · shapes	the method itself, and output others depend on	what good looks like for this work across teams	almost nothing within the discipline

THE TEST THAT KEEPS THE SCALE HONEST

Two experienced assessors, working independently and without discussing the person, should place them at the same level most of the time. If they do not, the anchors are not behavioural enough. Fix the anchors before you collect a single data point, because every downstream calculation inherits this noise and no amount of modelling removes it later.

quarter one is finished when

- ☐ between 30 and 50 roles are decomposed, each with 3 to 7 responsibilities and 8 to 12 skills
- ☐ every skill traces to at least one named task, and every responsibility to at least one KPI or a written note explaining why none exists
- ☐ the proficiency scale has behavioural anchors and has survived a two-assessor calibration test
- ☐ the line leader for each role has read the output and confirmed they recognise the job
- ☐ the decision the programme will change by quarter four is written down and has a named owner

quarter two: choosing a taxonomy you can maintain

the decision is not which taxonomy is best, it is which one you can still be running in three years.

QUARTER TWO AT A GLANCE

	what it is
entry criteria	role architecture complete for the selected roles, with skills named in your own language
the work	select a taxonomy, map your skill names to it, and commit to a refresh cadence
owners	a central team owns the taxonomy and the mapping; nobody else may add a skill
exit criteria	every skill in the architecture maps to exactly one taxonomy entry, with a documented refresh interval
common failure	selecting the largest taxonomy available and treating the purchase as the maintenance plan

A taxonomy is a maintenance commitment wearing the costume of a purchase. The selection decision should therefore be made on maintainability first and coverage second.

how fast the target moves

32% of the skills required by the average job changed over the three years to 2024, with roughly 75% turnover in the fastest-moving quarter of occupations
Lightcast, The Speed of Skill Change in the US, 2024

Lightcast's earlier work on fifteen million job postings found that 37 per cent of the top twenty skills requested for the average job had changed since 2016, and that around one in five was entirely new to the occupation.

Employer expectations point the same way. The World Economic Forum's Future of Jobs Report 2025 found that employers expect 39 per cent of key skills to change by 2030.

Take the conservative end of that range and assume drift of roughly ten per cent a year. A taxonomy refreshed every three years carries an average error near fifteen per cent, and a peak error near thirty per cent in the months before the refresh lands. No governance process survives being wrong about a third of its own content.

THE DRIFT BUDGET

Set a refresh interval per layer and treat missing it as a defect rather than a delay. A practical starting point: capability spine every 24 months, role profiles every 12, evidence and tooling every 3. Then report adherence as a metric. A model nobody refreshes is not a model, it is an archive.

the three options, honestly compared

TAXONOMY OPTIONS

option	strength	cost you will actually pay
purchased commercial taxonomy	large coverage, market benchmarking, refreshed by the vendor	annual licence, and a mapping exercise every time your role architecture changes
open standard such as ESCO or O*NET	free, stable, publicly documented, defensible to a regulator	coverage gaps in your sector, and refresh on the publisher's timetable rather than yours
inferred from your own data	reflects the work you actually do, no mapping layer	you own the entire maintenance burden, including deduplication and drift detection

The United States Department of Labor sponsors O*NET, which maintains around nine hundred occupational profiles covering more than fifty-five thousand jobs, with parts of every profile refreshed annually and fast-moving elements refreshed more often. A national labour agency has concluded that occupational description cannot be maintained as a periodic project. Enterprises with far fewer roles still attempt the three-year rewrite.

The European alternative is ESCO, the multilingual classification maintained by the European Commission, which covers skills, competences and occupations across all official EU languages. It is free, publicly documented and defensible to a works council or a regulator, which matters more than it appears once inferred data starts influencing decisions. Its limitation is the same as any public standard: it moves on the publisher's timetable and will lag your sector.

the taxonomy fitness test

score any candidate taxonomy out of ten. below seven, you will not be running it in three years.

01 refresh cadence

How often is it updated, and by whom? Score 2 for monthly or quarterly with a published changelog, 1 for annual, 0 for on request.

02 mapping cost

How much work is it to map your existing role architecture to it? Score 2 if a majority of your skills map one to one, 0 if most need judgement.

03 granularity fit

Does it distinguish at the level your decisions need? Score 2 if it separates skills your staffing decisions actually turn on, 0 if it lumps them.

04 exit cost

If you leave the vendor, what survives? Score 2 if your mappings and proficiency data are portable, 0 if the structure is proprietary and locked.

05 defensibility

Could you explain the classification to a regulator or a works council? Score 2 for a public standard or a documented method, 0 for a black box.

THE RULE THAT KEEPS IT CLEAN

One team owns the taxonomy and nobody else may add a skill. The moment three functions can each create entries, you have three taxonomies sharing a database, and the deduplication cost exceeds the original build.

mapping is a project with an end date, and it needs one

Once a taxonomy is selected, your role architecture has to be mapped onto it. This is the least interesting work in the whole programme and the most frequently underestimated. Give it a fixed end date and a rule for the hard cases.

1. Map the unambiguous cases first, in bulk. Most of your skills will map one to one and need no judgement.
2. For each remaining skill, decide once whether it is a naming difference or a genuine gap in the taxonomy. Naming differences get an alias. Genuine gaps get a local extension.
3. Cap local extensions at a stated share of the model, commonly around one in ten. Beyond that, you have chosen the wrong taxonomy and should say so rather than extend your way to a private one.
4. Record every mapping decision with a one-line rationale. In two years somebody will ask why two similar skills were treated differently, and the answer needs to exist.
5. Publish the mapping and freeze it for the quarter. A mapping that changes weekly makes every downstream comparison meaningless.

The cap on local extensions is the discipline that matters. Every extension is a piece of the model that no vendor will maintain for you, that no benchmark will cover, and that your team now owns

permanently.

Layer the model so that the parts moving at different speeds are separated. A capability spine of twelve to twenty durable capabilities changes every three to five years. Role profiles change annually. Evidence and tooling, the specific platforms and certifications, change quarterly. Most models rot because tooling-level content ends up in the spine, where a capability named after a vendor product is obsolete the moment procurement changes suppliers.

THREE LAYERS, THREE CLOCKS

layer	example	changes	who maintains it
capability spine	commercial negotiation, data interpretation, service design	every 3 to 5 years	the central team, deliberately and rarely
role profiles	the skills and proficiencies a named role requires	annually	the line leader, with a central review
evidence and tooling	a named platform, certification or method	quarterly	whoever owns the tool, automatically where possible

If your capability spine contains a vendor's product name, the spine will be wrong the next time procurement changes its mind.

quarter three: inference, validation and confidence thresholds

inferred skill data is useful long before it is accurate, provided you say out loud how accurate it is.

QUARTER THREE AT A GLANCE

	what it is
entry criteria	taxonomy selected and mapped; at least two systems that hold evidence of what people actually do
the work	infer skill profiles from work evidence, validate a sample, publish confidence thresholds
owners	people analytics owns inference and validation; line leaders own exceptions
exit criteria	a published confidence ladder, and a measured precision figure for at least one job family
common failure	shipping inferred profiles to managers without a stated confidence level, which burns trust once and permanently

Populating skill profiles by asking people is the most common approach and the least reliable. A self-rating measures self-perception. That is informative about interest, confidence and intent, and only weakly related to demonstrated capability.

Self-report is fine for routing people toward opportunities and for surfacing interests the organisation cannot otherwise see. It is not a sound input to a capability gap number that will be presented to an executive committee.

where the better signal already sits

- project and assignment records, which show what a person was actually staffed on and for how long
- internal requisitions, which show what hiring managers ask for when they are spending money
- completed assessments and certifications, which carry a date and an issuer
- goal and review text, which names the work rather than the title
- systems of work such as ticketing, code, CRM or case management, where the trace falls out of the job rather than being collected

PREFER TRACES TO SURVEYS

Evidence you have to ask for decays. Evidence that falls out of the work does not. Where both exist, weight the trace higher and use the survey to fill gaps.

source	signal strength	collection cost	the catch
completed assignments and project records	high	low, if the data already exists	records scope poorly; a person listed on a project may have done very little of it
assessments and certifications	high	medium to high	dated, so they age, and they measure the test rather than the work
systems of work: tickets, code, cases, CRM	high	low once integrated	coverage is uneven across functions, which biases any comparison
internal requisitions and their outcomes	medium	low	tells you what managers ask for, which is not always what the work needs
goal and review text	medium	low	quality varies with the writer, not with the person being described
self-assessment	low	low	measures confidence and interest; useful for routing, weak for claims
manager assessment	medium	medium	carries the manager's exposure to the person, not the person's full range

say how confident you are, in advance

the confidence ladder

publish this before any inferred profile reaches a manager, and let the decision follow the confidence rather than the other way round.

01 level 1, indicative

Single weak signal, such as one self-rating or one keyword match. Use for suggesting learning content and surfacing interests. Never show as a capability claim.

02 level 2, corroborated

Two independent signals agree, at least one from a system of work. Use for shortlisting suggestions and internal opportunity matching, always with the evidence visible.

03 level 3, evidenced

A dated assessment, certification or completed assignment at the required scope. Use for staffing decisions, successor slates and gap reporting to leadership.

04 level 4, verified

Level 3 plus independent confirmation by someone outside the person's reporting line. Reserve for decisions that are hard to reverse, such as promotion into a critical role.

The ladder does two things at once. It stops weak data carrying heavy decisions, and it stops the programme waiting for perfect data before delivering anything. Level 2 is enough to make a talent

marketplace useful. Only level 3 belongs in a capability gap presented to a board.

Write the mapping from rung to permitted use down and publish it in the same place as the data. The failure it prevents is specific and common: a level 1 profile built from self-ratings gets exported into a slide, the slide reaches an executive committee, and a headcount decision is taken on data that was never meant to leave a recommendation engine.

validate a sample, and publish the number

1. Draw a stratified sample of 150 to 300 inferred skill claims across job families and confidence levels.
2. Have a subject matter expert who does not know the inferred value assess each claim independently.
3. Compute precision, the share of inferred claims the expert confirms, separately for each confidence level.
4. Publish the precision figure alongside the data, in the interface where managers see it.
5. Repeat annually, and after any change to the inference model or the taxonomy.

Measure precision rather than recall in the first year. Precision asks whether the claims you make are true. Recall asks whether you found everything. A programme that overstates capability loses trust immediately, and trust is the constraint in year one. A programme that misses capability loses opportunity, which is recoverable and which better coverage will fix later.

A published precision figure of seventy per cent is more useful than an unstated one of ninety. Managers calibrate to a number they can see. They cannot calibrate to a claim of accuracy with no evidence behind it, so they discount it entirely, which is the rational response.

what to tell employees, before they ask

Inferring capability from systems of work touches people data, and in most jurisdictions it engages employment law, data protection obligations and, in much of Europe, a consultation requirement. Handle this at the start rather than after the first complaint.

- state which systems are used as evidence, in plain language, before inference runs
- give every person the ability to see their own inferred profile and to challenge any entry
- record a challenge as a data point rather than an exception, because a pattern of challenges in one job family is a finding about the model
- never let an inferred profile be the sole basis for a decision that affects someone's employment
- where a works council or employee representative body exists, consult before the first inference run rather than before the first decision

THE PRACTICAL REASON, NOT ONLY THE LEGAL ONE

An employee who can see and correct their profile becomes a source of validation you did not have to fund. An employee who discovers an inferred profile they were never shown becomes the reason the programme is paused.

THE TRUST ASYMMETRY

Trust in inferred data is lost faster than it is built. One confident wrong recommendation to a senior manager costs more than fifty accurate ones earn. This is the argument for shipping at level 2 with visible evidence rather than at level 1 with confidence.

before any inferred profile reaches a manager

- ☐ the confidence ladder is published, and every profile displays its rung
- ☐ measured precision exists for at least one job family and is visible next to the data
- ☐ the evidence behind every level 2 claim can be inspected by the person it describes
- ☐ employees have been told which systems are used as evidence, and how to challenge an entry
- ☐ no decision that affects employment can be taken on a level 1 or level 2 claim alone

quarter four: skills-based staffing and the first real decisions

this is the quarter where the resistance appears, because it is the first quarter in which anything is at stake.

QUARTER FOUR AT A GLANCE

	what it is
entry criteria	one job family with validated profiles at confidence level 2 or better, and a named hiring leader willing to try
the work	rewrite requisitions for one job family as skills-based, and run them for a full cycle
owners	the hiring leader owns the decision; talent acquisition owns the process change
exit criteria	a measured difference in shortlist composition, and a written account of what the hiring managers rejected
common failure	piloting across many families at once, which produces weak signal everywhere and a committed sponsor nowhere

Everything before this quarter was reversible and cost nobody anything. Skills-based staffing is the first point at which a manager is asked to consider a candidate they would not otherwise have seen. That is where a skills programme either becomes real or becomes a dashboard.

pick the family for the right reason

The instinct is to pilot where the data is cleanest. Resist it. Pilot where the pain is loudest. A family with clean data and no hiring problem produces a technically successful pilot that nobody defends at budget time. A family with messy data and a leader who has been complaining about time to fill for two years produces an advocate, and an advocate is what carries the programme into year two.

the staffing switch

five steps, run on one job family, in one quarter.

01 pick the family

Choose one with frequent hiring, a real internal candidate pool, and a leader who has complained about time to fill. Volume gives you signal and the complaint gives you a sponsor.

02 rewrite the requirement

Convert the requisition from a title and a years-of-experience figure into the 8 to 12 skills with required proficiency. Keep any genuine legal or licensing requirement, and delete proxies such as degree class and prior title.

03 run a parallel shortlist

For the first cycle, produce both the traditional shortlist and the skills-based one. Do not tell the hiring manager which is which.

04 record every rejection

When a manager rejects a skills-based candidate, record the stated reason. This is the highest-value data the whole programme produces, because it names what the model is missing.

05 change one thing

At the end of the cycle, change the single input that explains most rejections. Resist the urge to rebuild the model.

The parallel shortlist is what turns an argument into evidence. Without it, the conversation reduces to whether the hiring manager trusts the system, which is a conversation nobody wins.

WHAT TO MEASURE ACROSS THE PILOT CYCLE

measure	traditional shortlist	skills-based shortlist	what a difference means
internal candidates surfaced	baseline	compare	the marketplace is finding people the process was missing
candidates progressing past first screen	baseline	compare	the matches are real, not merely numerous
time from requisition to offer	baseline	compare	the practical case for the finance director
hiring manager rejection reasons	record	record	the diagnostic that improves the model
performance at six months	follow up	follow up	the only measure that settles the argument, and it arrives late

The last row is the one to commit to in writing at the start, because it is the one everyone forgets. A skills-based hire who succeeds at six months is the single most persuasive artefact the programme will ever produce, and it costs almost nothing to capture if you decide to before you need it.

what the rejections will tell you

- a missing skill nobody named in the role session, which means the architecture needs one addition
- a proficiency level set too low, which is the most common finding and the easiest to correct
- a genuine contextual requirement such as a regulated market or a language, which belongs in the profile as a constraint
- a preference that cannot be stated in skill terms, which is the finding worth escalating, because it is often where bias lives

EXPECT THIS ONE

The most frequent objection is that the internal candidate lacks industry experience. Ask which specific task they could not perform. If nobody can name one, the objection is a proxy, and the requisition should not carry it.

the conversation with the hiring manager

This conversation decides the pilot. It goes better when it is framed as a request for their judgement rather than a request for their compliance.

1. Open with what they lose today. Time to fill, candidates who withdraw, the internal person who left because they never saw the opening.
2. Ask them to name the three capabilities that separate a strong performer from an adequate one in their team. Use their words in the requisition.
3. Show both shortlists without labels, and ask which candidates they want to meet.
4. When they reject a skills-based candidate, ask which specific task the person could not perform. Write down the answer without arguing with it.
5. Come back after the cycle with what changed in the model because of their feedback. This is what makes them a participant rather than a test subject.

The World Economic Forum's Future of Jobs Report 2025 puts the scale of the retraining task plainly. Of every hundred workers, 59 will need training by 2030. Twenty-nine can be upskilled within their current role, 19 will need reskilling and redeployment elsewhere in the business, and 11 are unlikely to receive the training they need.

That last group is the one a skills programme either accounts for or abandons by omission. Quarter four is where you find out which of those you are running.

A skills programme becomes real on the day a manager hires someone they would not have shortlisted, and it works on the day that person succeeds.

pilot readiness

- ☐ one job family, one hiring leader, and a written statement of what would count as success
- ☐ profiles for that family validated at confidence level 2 or better, with a measured precision figure
- ☐ the requisition rewritten as skills with required proficiency, with proxies removed and genuine legal requirements kept
- ☐ a mechanism to capture every rejection reason as structured data, not as an email thread
- ☐ agreement, in advance, to follow up on performance at six months

rewiring job architecture, pay bands and promotion

the plumbing behind titles is what quietly reverses a skills programme in year two.

A skills programme can survive quarter four and still be undone by three systems that were never in scope: job architecture, pay and promotion. Each of them is anchored to titles, and each of them carries more institutional weight than the skills model does.

job architecture: keep the grade, change what it means

The instinct is to replace the job architecture. Resist it. Grades carry pay, benefits, approval authority, and in several jurisdictions statutory and works council implications. Replacing them is a multi-year programme with its own failure modes.

The workable move is to keep the grade structure and change what determines placement within it. A grade stops being a statement about title and seniority and becomes a statement about the scope and proficiency a role requires.

In practice this is a small number of edits to documents that already exist, which is the reason it works.

- the grade definition gains a statement of the capability depth expected at that grade, alongside the existing scope language
- the job description template gains a skills block with required proficiency, and loses the years-of-experience line
- the requisition form requires skills before it will accept a title
- the internal transfer process references the skill profile rather than the prior title

None of those four changes requires renegotiating pay, restructuring grades, or consulting on a new architecture. All four change what the organisation asks for by default, which over eighteen months does more than a redesign would.

pay: two mechanisms, and only one of them is safe to start with

TWO WAYS TO PAY FOR SKILLS

mechanism	how it works	risk
skill premium	a defined, time-limited supplement for holding a scarce and verified capability	low, because it is reversible and sits outside base pay
skills-based base pay	base pay determined by the skill set held rather than the role occupied	high, because it is difficult to reverse and interacts with equal pay obligations

Start with premiums attached to verified, dated capability at confidence level 3 or above. They are reversible, they are explainable, and they test whether your assessment is defensible before anything permanent depends on it.

A PAY TRANSPARENCY INTERACTION

In jurisdictions implementing pay transparency obligations, any pay differential must be justifiable on objective criteria. A skill premium backed by a dated assessment is easier to defend than a discretionary allowance. A premium backed by a self-rating is harder.

European employers have a fixed date to work toward. The EU Pay Transparency Directive was adopted in 2023 and member states must transpose it into national law by 7 June 2026. It requires pay structures to rest on objective, gender-neutral criteria and gives workers the right to information about pay levels for work of equal value. A skills model with dated, documented assessment is an asset under that regime. An undocumented one becomes a liability the first time a differential is questioned.

promotion: separate the two things a promotion currently does

A promotion in most organisations bundles two distinct claims: that a person has grown, and that the organisation has a larger job for them. Bundling them means growth is only recognised when a vacancy exists, which is why capable people leave in flat structures.

- recognise proficiency growth through progression within a grade, tied to verified capability, available at any time
- reserve promotion for a genuine change in scope, where the role itself is larger
- publish both routes, so that a person can see which one they are on

This separation is the single change that does most to make a skills model believable to employees, because it gives the model a consequence they can see without waiting for someone to leave.

consultation is not an afterthought

Anything touching pay, grading or promotion criteria engages employee representative structures in much of Europe and in several Asian jurisdictions. Works councils commonly hold information and consultation rights over changes to remuneration systems and over the introduction of systems capable of monitoring performance or behaviour, which inferred skill data can reasonably be argued to be.

SEQUENCE THIS DELIBERATELY

Consult on the method before the data exists, not on the results after they do. A works council asked to review a documented assessment approach is reviewing a proposal. The same council shown a finished ranking of its members is being presented with a fait accompli, and will treat it as one.

Skills change what the organisation values. Pay and promotion are how an organisation states what it values. Leave them untouched and the statement contradicts the model.

before you touch pay

- ☐ the capability behind any premium is assessed at confidence level 3 or above, with a date
- ☐ the premium has a stated review interval and a defined route to removal
- ☐ legal and reward have reviewed the mechanism against local equal pay and transparency obligations
- ☐ the assessment method behind the premium would survive being shown to the person who did not receive it
- ☐ you can state, in one sentence, what the organisation gets in return for the premium

the metrics that show it is working, and the ones to drop

most skills dashboards report the size of the artefact, which is the one thing that does not matter.

A skills programme is easy to measure badly, because the artefact produces attractive numbers. Skill counts rise. Profile completion rises. Neither tells you whether anything changed.

The test for any metric on a skills programme is whether a bad result would cause someone to do something different. Most skills metrics fail it. A falling skill count triggers nothing. A falling decision coverage figure triggers a conversation with a hiring leader, which is the conversation the programme exists to have.

drop these

METRICS TO STOP REPORTING

metric	why it misleads
number of skills in the taxonomy	grows without limit and correlates with maintenance burden, not with capability
profile completion percentage	measures data entry, and rises fastest when the data is self-reported and unvalidated
number of skills per employee	rewards inflation, and a longer profile is harder to act on rather than easier
learning hours or course completions	measures consumption, and has no established link to capability at the point of use
taxonomy coverage against an external benchmark	measures similarity to a vendor's model, not fitness for your decisions

keep these

METRICS WORTH AN EXECUTIVE REVIEW

metric	what it tells you	healthy direction
decision coverage	share of open roles in scope where a skills-based shortlist was produced	rising toward the families you have architected
inference precision	share of inferred skill claims confirmed on validation, by confidence level	stable and published, not necessarily high
internal fill rate	share of open roles filled internally, for architected families	rising, compared against families not yet in scope
drift budget adherence	share of role profiles refreshed within their stated interval	above 90 per cent, or the interval is wrong
rejection reason mix	why hiring managers reject skills-based candidates	shifting from missing skills toward genuine constraints

Report the first four as a set rather than individually. Rising decision coverage with flat internal fill means the shortlists are being produced and ignored. High precision with low decision coverage means the data is trusted and unused. Each pair diagnoses a different problem.

the twelve-month scoreboard

WHAT TO BE ABLE TO REPORT AT EACH QUARTER END

quarter	the number	the artefact
Q1	roles decomposed, and the share whose line leader has confirmed them	a role architecture for 30 to 50 roles, and one named decision
Q2	share of skills mapped, and the share requiring a local extension	a selected taxonomy, a frozen mapping and a published refresh cadence
Q3	measured precision by confidence level, for at least one job family	a published confidence ladder and a validation result
Q4	decision coverage, and internal fill rate for the pilot family	a completed hiring cycle with recorded rejection reasons

the quarterly review question set

four questions, asked in the same order every quarter, of the same three people.

01 which decision changed

Name one staffing, development or pay decision that came out differently this quarter because of skills data. If there is none, the programme is in the artefact stall.

02 what did we get wrong

Name the rejection reason that appeared most often, and what changed in response. A quarter with no corrections means nobody is using the output.

03 what is stale

Report drift budget adherence by family. Anything below the threshold gets an owner and a date, not an explanation.

04 what are we not measuring

Name one capability the business depends on that the model does not yet describe. This is the queue for next quarter's architecture work.

ONE NUMBER FOR THE BOARD

If you can report only one figure, report decision coverage. It is the only metric that fails when the programme becomes an artefact, which is the failure mode this guide exists to prevent.

what year two looks like if year one worked

A successful first year does not end with a finished model. It ends with a maintained one, a hiring leader who will say in public that it helped, and a queue of the next decisions to rewire. Year two is mostly extension, and extension is far easier to fund than initiation.

1. Extend the role architecture to the next thirty roles, using the method that already works rather than revisiting it.
2. Add a second decision: development planning, succession slates or project staffing. One at a time, each with the same rejection-reason discipline.
3. Move the pilot family from confidence level 2 to level 3 by adding assessment where the stakes justify the cost.
4. Introduce skill premiums for two or three genuinely scarce capabilities, each with a review date and a route to removal.
5. Report the twelve-month scoreboard against the first year's baseline, including what you stopped doing.

A skills model that describes forty roles accurately, refreshes monthly, and changes one class of decision is worth more than one describing four thousand roles that nothing consumes. Coverage can always be extended later. A programme that never touched a decision rarely gets a second year.



Build the vocabulary if you like. But fund the decision, or the vocabulary is all you will have.

REFERENCES

sources

every figure in this document traces to one of the following. links were live at the date of publication.

- 1 Deloitte Insights, The skills-based organization: A new operating model for work and the workforce, 2022
<https://www.deloitte.com/us/en/insights/topics/talent/organizational-skill-based-hiring.html>
- 2 Harvard Business School and Burning Glass Institute, Skills-Based Hiring: The Long Road from Pronouncements to Practice, 2024
<https://www.burningglassinstitute.org/research/skills-based-hiring-2024>
- 3 Lightcast, The Speed of Skill Change in the US, 2024
<https://lightcast.io/resources/research/speed-of-skill-change>
- 4 Lightcast, Burning Glass Institute and BCG, Shifting Skills, Moving Targets, and Remaking the Workforce, 2022
<https://lightcast.io/resources/research/shifting-skills>
- 5 World Economic Forum, Future of Jobs Report 2025
<https://www.weforum.org/publications/the-future-of-jobs-report-2025/>
- 6 O*NET Resource Center, Data Collection Overview, 2025
<https://www.onetcenter.org/dataCollection.html>
- 7 O*NET Resource Center, Content Model, 2025
<https://www.onetcenter.org/content.html>
- 8 European Commission, ESCO classification of skills and competences, 2025
https://esco.ec.europa.eu/en/classification/skill_main
- 9 Deloitte Insights, Global Human Capital Trends: A skills-based model for work, 2023
<https://www.deloitte.com/us/en/insights/topics/talent/human-capital-trends/2023/skills-based-model-end-of-jobs.html>
- 10 Directive (EU) 2023/970, EU Pay Transparency Directive, 2023
<https://eur-lex.europa.eu/eli/dir/2023/970/oj>
- 11 Lightcast, Open Skills library and taxonomy documentation, 2025
<https://lightcast.io/open-skills>
- 12 SkillsFuture Singapore, Skills Framework, 2025
<https://www.skillsfuture.gov.sg/skills-framework>

about entomo

entomo is an agentic workforce intelligence platform for performance, skills, and talent decisions, used by more than 30 million people worldwide. its agents reason on axon, a living model of work that connects what every role requires with what every person brings, benchmarked against the market and refreshed every month.

the practice behind this document is the same one behind the platform: HR process depth first, intelligence second.

see it against your own hardest role

entomo.co

